



北京交通大学

CPU Scheduling





Basic Concepts

Scheduling

- Long-Term Scheduling

- Determines which programs are admitted to the system for processing
- Controls the degree of multiprogramming

- Medium-Term Scheduling

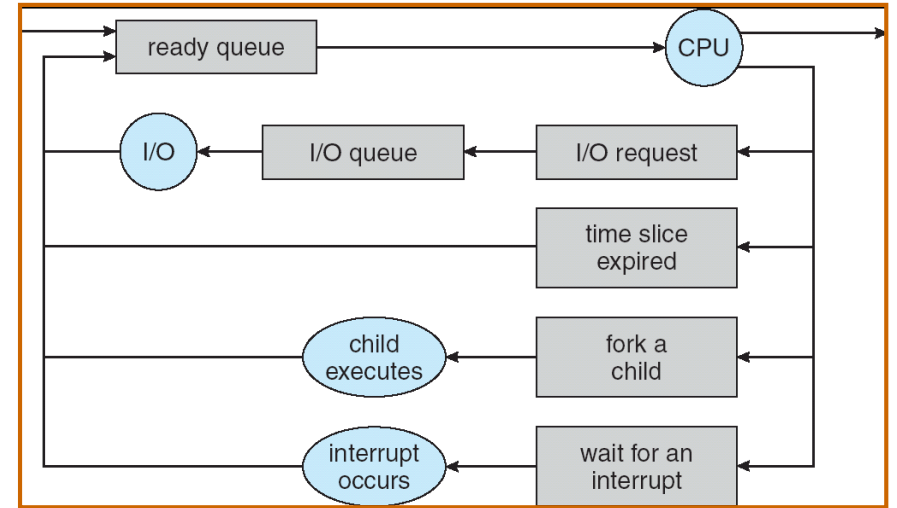
- Part of the swapping function
- Based on the need to manage the degree of multiprogramming

- Short-Term Scheduling

- Known as the dispatcher
- Executes most frequently

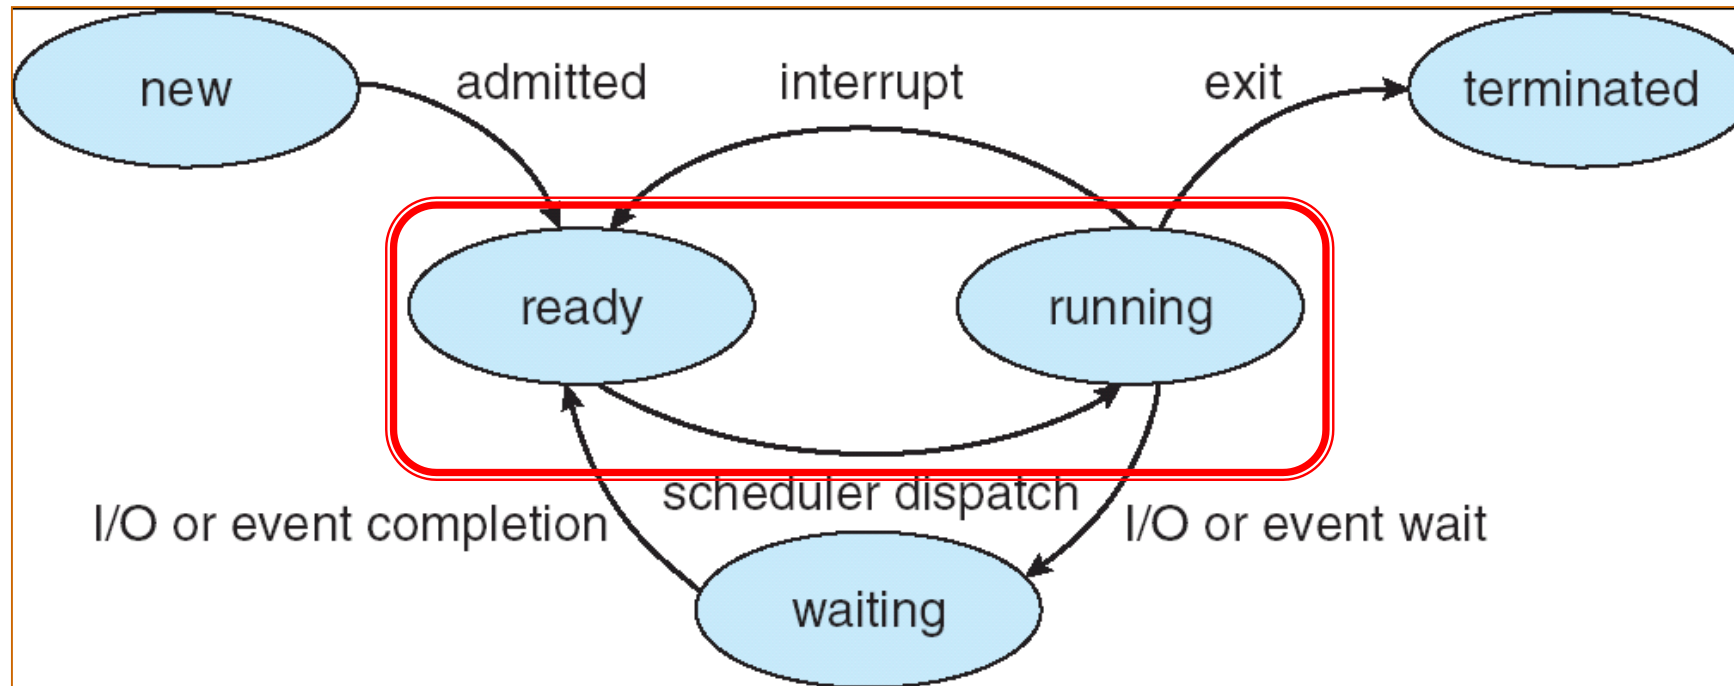
CPU Scheduling

- Earlier, we talked about the life-cycle of a process



- Question: How is OS to decide which of several tasks to take off a queue?
 - Obvious queue to worry about is ready queue
 - Others can be scheduled as well, however

Diagram of Process State




CPU Scheduler

- **Selects** from the processes in memory that are ready to execute, and **allocates** the CPU to one of them
- CPU scheduling decisions may take place when a process:
 - 1. Switches from running to waiting state
 - 2. Switches from running to ready state
 - 3. Switches from waiting to ready
 - 4. Terminates
- Scheduling under 1 and 4 (no choice in terms of scheduling) is **nonpreemptive**
- All other scheduling is **preemptive**

Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler
- **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running



Scheduling Assumptions

- CPU scheduling is a big area of research in early 70's
- Many implicit assumptions for CPU scheduling:
 - One program per user
 - One thread per program
 - Programs are independent

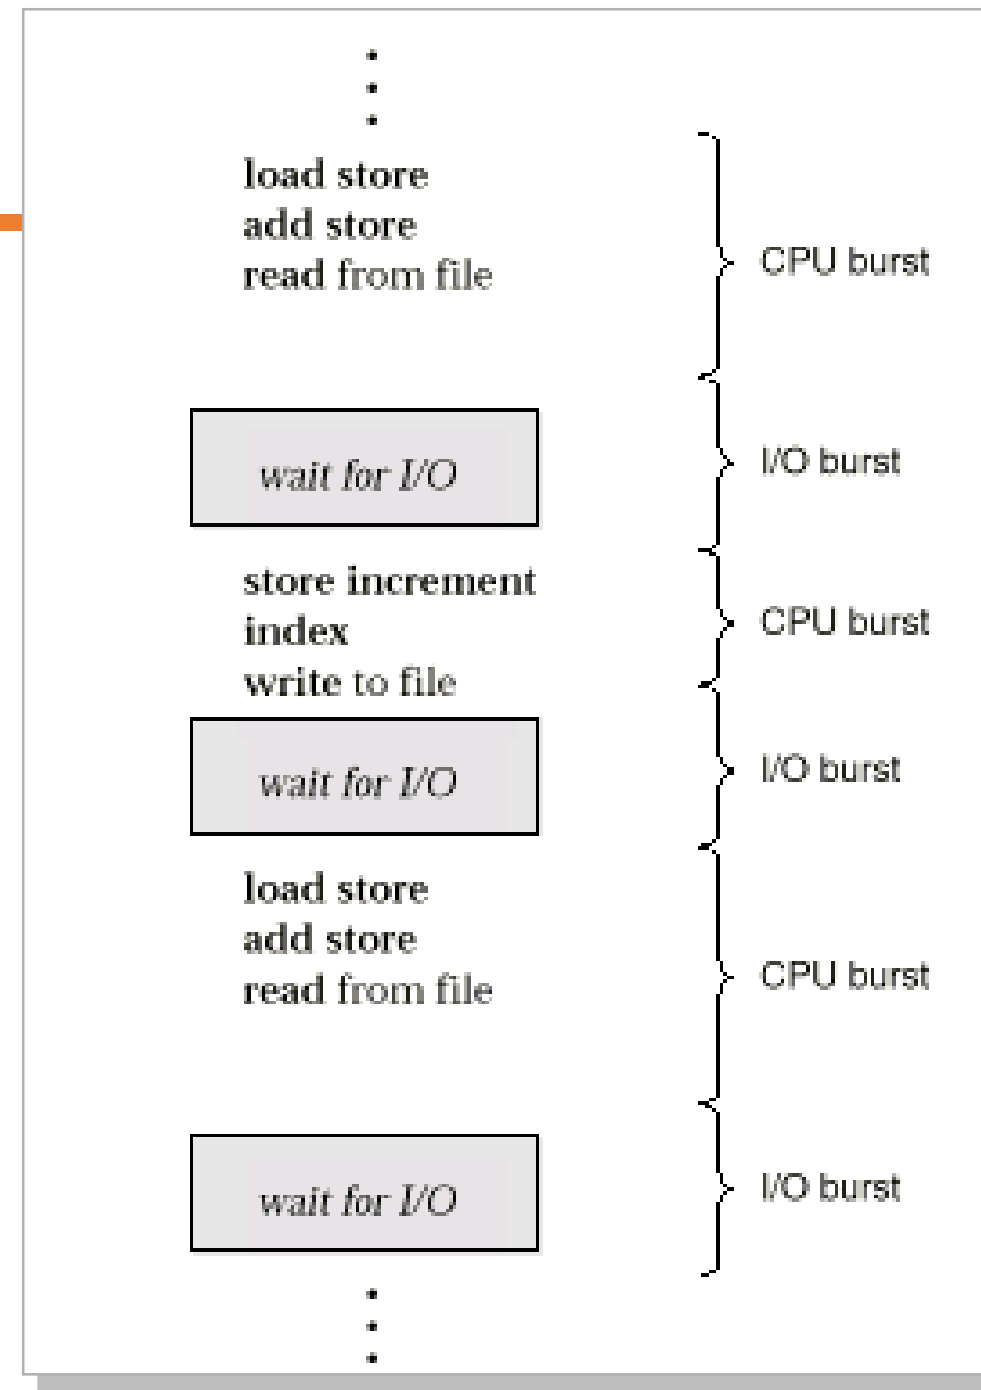


Scheduling Assumptions

- Clearly, these are unrealistic but they simplify the problem so it can be solved
 - For instance: is “fair” about fairness among users or programs?
 - If I run one compilation job and you run five, you get five times as much CPU on many operating systems
- The **high-level goal**: Dole out CPU time to optimize some desired parameters of system


Basic Concepts

- **CPU-I/O burst cycle:** Process execution consists of a cycle of CPU execution and I/O wait (i.e., **CPU burst** and **I/O burst**).
- Execution model: programs alternate between bursts of CPU and I/O



CPU Burst Distribution

- Histogram of CPU-burst Times
(exponential or hyper-exponential)





Scheduling Criteria

Scheduling Policy Goals/Criteria

- Minimize Response Time
- Maximize Throughput
- Fairness

Scheduling Policy Goals/Criteria

- **Minimize Response Time**
 - Response time (of a request)
 - submission ~ the first response is produced
 - Response time is what the user sees:
 - Time to echo a keystroke in editor
- Maximize Throughput
- Fairness

Scheduling Policy Goals/Criteria

- Minimize Response Time
- Maximize Throughput
 - Throughput
 - number of completed processes per time unit
 - Maximize operations (or jobs) per second
- Fairness

Scheduling Policy Goals/Criteria

- Minimize Response Time
- Maximize Throughput
- Fairness
 - Share CPU among users in some equitable way

Scheduling Policy Goals/Criteria

- CPU utilization
 - theoretically: 0%~100%
 - real systems: 40% (light)~90% (heavy)
- Turnaround time (of a process)
 - submission ~ completion
- Waiting time (of a process)
 - total waiting time in the ready queue



Scheduling Algorithms

Scheduling Algorithms

- First-Come-First-Served (FCFS) Scheduling
- Round-Robin (RR) Scheduling
- Shortest-Job-First (SJF) Scheduling
- Shortest-Remaining-Time-First (SRTF) Scheduling
- Priority Scheduling
- Multilevel Queue Scheduling
- Multilevel Feedback Queue Scheduling

First-Come, First-Served (FCFS) Scheduling

- First-Come, First-Served (FCFS)
 - Also “First In, First Out” (FIFO) or “Run until done”
 - In early systems, FCFS meant one program scheduled until done (including I/O)
 - Now, means keep CPU until process blocks






FCFS Scheduling

- Each process joins the Ready queue
- When the current process ceases to execute, the **oldest** process in the ready queue is selected




FCFS Scheduling



FCFS Scheduling

- Example: 4 processes arrived almost simultaneously



- Waiting time:

$$P_1=0; P_2=3; P_3=7; P_4=19$$

- Average waiting time:

$$(0+3+7+19)/4 = 7.25$$


- Average completion time:

$$(3+7+19+35)/4=16$$



FCFS Scheduling

- Example: 4 processes arrived almost simultaneously




- Waiting time: $P_2=0$; $P_3=4$; $P_1=16$; $P_4=19$
- Average waiting time: $(0+4+16+19) / 4 = 9.75$
- Average completion time: $(4+16+19+35) / 4 = 18.5$



FCFS Scheduling

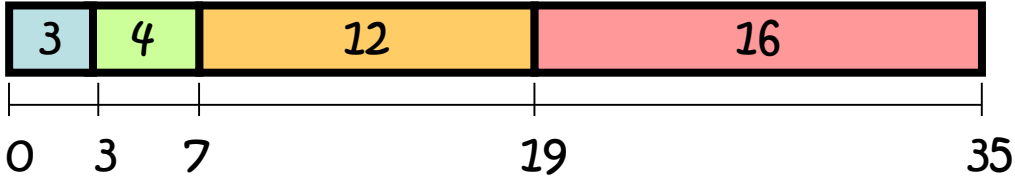
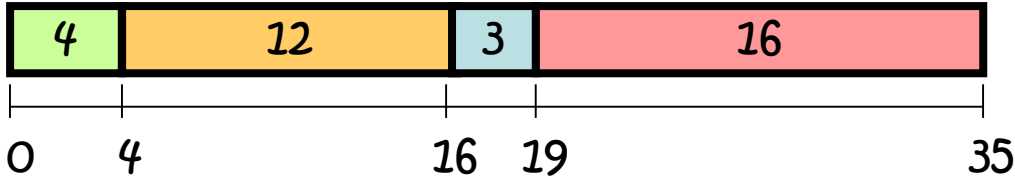
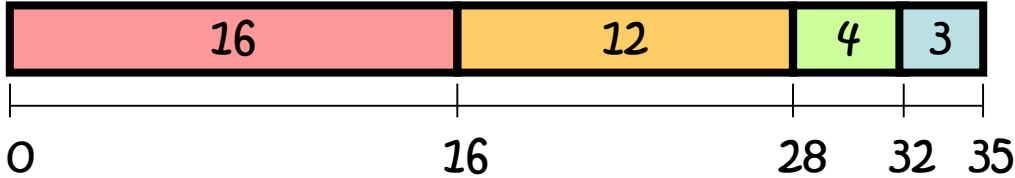
- Example: 4 processes arrived almost simultaneously



- Waiting time: $P_4=0$; $P_3=16$; $P_2=28$; $P_1=32$
- Average waiting time: $(0+16+28+32) / 4 = 19$
- Average completion time: $(16+28+32+35) / 4 = 27.75$



FCFS Scheduling

Scheduling Example	Average waiting time	Average completion time
 <p>Timeline: 0, 3, 7, 19, 35</p>	7.25	16
 <p>Timeline: 0, 4, 16, 19, 35</p>	9.75	18.5
 <p>Timeline: 0, 16, 28, 32, 35</p>	19	27.75

FCFS Scheduling

- Simplest(+), non-preemptive, often having a long average waiting time(-)
- A short process may have to wait a very long time before it can execute(-)
- Favors CPU-bound processes
 - I/O processes have to wait until CPU-bound process completes

FCFS Scheduling

- FCFS Scheme: **Potentially bad for short jobs!**
 - Depends on submit order
 - If you are first in line at supermarket with milk, you don't care who is behind you, on the other hand...



Round Robin (RR)

- Round Robin Scheme
- The modern use of the term dates from the 17th Century French ruban rond (round ribbon)




Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum**), usually 10-100 milliseconds
- After quantum expires, the process is **preempted** and **added to the end of the ready queue**.
- n processes in ready queue, time quantum is $q \Rightarrow$
 - Each process gets $1/n$ of the CPU time
 - No process waits more than $(n-1)q$ time units

-
- Diagram illustrating the hold times for four processes:
- P_2 (green box) has a hold time of 4.
 - P_3 (orange box) has a hold time of 12.
 - P_1 (light blue box) has a hold time of 3.
 - P_4 (pink box) has a hold time of 16.

-
- A diagram of a 1D lattice with 20 sites. The first 10 sites are occupied by particles (colored squares). The last 10 sites are empty. A blue line with arrows indicates a path from the first site to the last site, passing through the occupied sites.

Round Robin (RR)




Example of RR with Time Quantum = 20

- Example: 4 processes arrived almost simultaneously




- The Gantt chart is:




Example of RR with Time Quantum = 20

- Example: 4 processes arrived almost simultaneously



- The Gantt chart is:




• Waiting time


- $P_1 = (68 - 20) + (112 - 88) = 72$; $P_2 = (20 - 0) = 20$; $P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$; $P_4 = (48 - 0) + (108 - 68) = 88$
- Average waiting time** = $(72 + 20 + 85 + 88) / 4 = 66.25$
- Average completion time** = $(125 + 28 + 153 + 112) / 4 = 104.5$

Round-Robin Pros and Cons

- Better for short jobs, Fair (+)
- Context-switching time adds up for long jobs (-)




Time Quantum & Context Switch Time




Round-Robin Discussion

- How do you choose **time slice**?
 - What if too big?
 - Response time suffers
 - What if infinite (∞)?
 - Get back FCFS
 - What if time slice too small?
 - Throughput suffers!
 - Must be large with respect to context switch, otherwise overhead is too high (all overhead)



Round-Robin Discussion

- Actual choices of time slice:
 - Initially, the UNIX time slice is 1s:
 - Worked ok when UNIX was used by one or two people.
 - What if three compilations going on? 3 seconds to echo each keystroke!
 - In practice, need to balance short-job performance and long-job throughput:
 - Typical time slice today is between 10ms – 100ms
 - Typical context-switching overhead is 0.1ms – 1ms
 - Roughly 1% overhead due to context-switching



Comparisons between FCFS and RR

- Assuming zero-cost context-switching time, is RR always better than FCFS?
- Simple example:
 - 10 jobs, each take 100s of CPU time
 - RR scheduler quantum of 1s
 - All jobs start at the same time
- Completion Time:
 - Both RR and FCFS finish at the same time
 - Average completion time is much worse under RR!
 - Bad when all jobs same length

Job #	FCFS	RR
1	100	991
2	200	992
...
9	900	999
10	1000	1000


Example with Different Time Quantum

Best FCFS: $P_2[8]$ $P_4[24]$ $P_1[53]$ $P_3[68]$

0 8 32 85 153

	Quantum	P_1	P_2	P_3	P_4	Average
Waiting Time	Best FCFS	32	0	85	8	$31\frac{1}{4}$
	Q = 1	84	22	85	57	62
	Q = 5	82	20	85	58	$61\frac{1}{4}$
	Q = 8	80	8	85	56	$57\frac{1}{4}$
	Q = 10	82	10	85	68	$61\frac{1}{4}$
	Q = 20	72	20	85	88	$66\frac{1}{4}$
	Worst FCFS	68	145	0	121	$83\frac{1}{2}$
Completion Time	Best FCFS	85	8	153	32	$69\frac{1}{2}$
	Q = 1	137	30	153	81	$100\frac{1}{2}$
	Q = 5	135	28	153	82	$99\frac{1}{2}$
	Q = 8	133	16	153	80	$95\frac{1}{2}$
	Q = 10	135	18	153	92	$99\frac{1}{2}$
	Q = 20	125	28	153	112	$104\frac{1}{2}$
	Worst FCFS	121	153	68	145	$121\frac{3}{4}$

Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

5 3 1 5 1 2

15+

8+


9+

17 = 49

49/4 = 12.25

Discussion

- Average waiting time = $(3a+2b+c)/4$
- Average completion time = $(4a+3b+2c+d)/4$



What if we knew the future?

- Could we always mirror best FCFS?
- Shortest Job First (SJF):
 - **Idea** is to get short jobs out of the system
 - Run whatever job has the **least amount of computation** to do
 - Sometimes called “Shortest Time to Completion First” (STCF)
- Associate with each process the length of its next CPU burst.



What if we know the future?

- Shortest Remaining Time First (SRTF):
 - Preemptive version of SJF
 - if a job arrives and has a shorter time to completion than the remaining time on the current job, immediately preempt CPU
 - Sometimes called “Shortest Remaining Time to Completion First” (SRTCF)




Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P ₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4

Example of Non-Preemptive SJF 北京交通大学

Process	Arrival Time	Burst Time
P₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4

- SJF (non-preemptive)



- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$


Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P ₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4

Example of Preemptive SJF

Process	Arrival Time	Burst Time
P ₁	0.0	7 5
P ₂	2.0	4 2
P ₃	4.0	1
P ₄	5.0	4

- SRTF (preemptive)



- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$

Discussion

- SJF/SRTF are the **best** you can do at minimizing average waiting time
 - Provably optimal (SJF among non-preemptive, SRTF among preemptive)
 - Since SRTF is always at least as good as SJF, focus on SRTF

Discussion

- Comparison of SRTF with FCFS and RR
 - What if all jobs the same length?
 - SRTF becomes the same as FCFS (i.e. FCFS is best can do if all jobs the same length)
 - What if jobs have varying length?
 - SRTF (and RR): short jobs not stuck behind long ones

SRTF Further discussion

- Starvation
 - SRTF can lead to **starvation** if many small jobs!
 - Short-running I/O bound jobs stay near top
 - CPU bound jobs drop like a rock
 - Large jobs may never get to run



SRTF Further discussion

- Bottom line, can't really know how long job will take
 - However, can use SRTF as a **yardstick** for measuring other policies
 - Optimal, so can't do any better
- SRTF Pros & Cons
 - Optimal (average waiting time) (+)
 - Hard to predict future (-)
 - Unfair (-)



SRTF Further discussion

- Somehow need to predict future
 - How can we do this?
 - Some systems **ask** the user
 - When you submit a job, have to say how long it will take
 - To stop cheating, system kills job if takes too long
- But: users have trouble predicting runtime of their jobs



SRTF Further discussion

- **Difficulty:** no way to know length of the next CPU burst. Thus, it cannot be implemented in short-term scheduler
- **Approximate SJF:** the next burst can be predicted as an exponential average of the measured length of previous CPU bursts


Determining Length of Next CPU Burst

- Can only **estimate** the length
- Using the length of previous CPU bursts by exponential averaging



1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} is the predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define: $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$.

Prediction of the Length of Next CPU Burst



$$\tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n.$$

CPU burst t_n	$t_0 =$	6	4	6	4	13	13	13	...
guess τ_n	$\tau_0 =$	10	8	6	6	5	9	11	...
Diff.		4	4	0	2	-8	-4	-2	...

Examples of Exponential Averaging

- $\alpha = 0$

- $\tau_{n+1} = \tau_n$

- Recent history does not count

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n.$$

- $\alpha = 1$

- $\tau_{n+1} = t_n$

- Only the actual last CPU burst counts

Examples of Exponential Averaging

- If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n.$$

- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor

Scheduling algorithms

- First-Come-First-Served (FCFS) Scheduling
- Round-Robin (RR) Scheduling
- Shortest-Job-First (SJF) Scheduling
- Shortest-Remaining-Time-First (SRTF) Scheduling
- Priority Scheduling
- Multilevel Queue Scheduling
- Multilevel Feedback Queue Scheduling

Priority scheduling


- A **priority number** (integer) is associated with each process
- The CPU is allocated to the process with the highest priority
 - Preemptive
 - Non-preemptive
- SJF?
 - is a priority scheduling where priority is the predicted next CPU burst time (smallest integer \equiv highest priority)

An Example

Example:

5 processes arrived almost simultaneously

Process	Burst Time	Priority
P1	10	3
P2	1	1
P3	2	4
P4	1	5
P5	5	2



Average waiting time = 8.2

An Example

Example:

5 processes arrived
almost simultaneously

Process	Burst Time	Priority
P1	10	3
P2	1	1
P3	2	4
P4	1	5
P5	5	2

Priority Scheduling

- Problem: **Starvation** – low priority processes may never execute
 - At MIT, there was a job submitted in 1967 that had not be run in 1973.
- Solution: **Aging** – as time progresses increase the priority of the process

Static priority scheduling

- Static priority scheduling
 - Prioritizing a process when it is created and keeping the priority unchanged until the process is finished.
- How to prioritize?
 - The type of process
 - System process › user process, interactive process › batch process
 - The need for system resources
 - Higher priority for processes with less CPU and memory requirements
 - User requirements
 - User level & payment, e.g. military computers, commercial computers

Dynamic priority scheduling

- **Dynamic priority scheduling**
 - The priority given to the process when it is created can be **dynamically** changed during process operation for better scheduling performance
- To prevent “starvation”, priority is adjusted based on running time and waiting time
 - A process lowers its priority every time it executes.
 - In the ready queue, the priority of the process is gradually increased as the waiting time increases.

Highest response ratio next

- **Highest Response Ratio Next (HRRN)** scheduling
 - a non-preemptive discipline, similar to SJF
 - the priority of each job is dependent on its estimated run time, and also the amount of time it has spent waiting.

$$\begin{aligned} \text{Priority} &= \frac{\text{estimated run time} + \text{waiting time}}{\text{waiting time}} \\ &= 1 + \frac{\text{estimated run time}}{\text{waiting time}} \end{aligned}$$

Highest response ratio next

- Jobs gain **higher priority** the longer they wait, which prevents indefinite postponement (process starvation).
- In fact, the jobs that have spent a long time waiting compete against those which are estimated to have a short running time.

$$\begin{aligned} \text{Priority} &= \frac{\text{estimated run time} + \text{waiting time}}{\text{waiting time}} \\ &= 1 + \frac{\text{estimated run time}}{\text{waiting time}} \end{aligned}$$

Example 1

JOB	ARRIVAL TIME	RUNNING TIME
0	10: 00	24 minutes
1	10: 10	60 minutes
2	10: 15	36 minutes
3	10: 20	12 minutes

- Assume that all jobs are completely CPU bound.
- For each of the scheduling algorithms, determine the **scheduling order** and the **average turnaround time**
 - FCFS
 - SJF

Example 1

JOB	ARRIVAL TIME	RUNNING TIME
0	10: 00	24 minutes
1	10: 10	60 minutes
2	10: 15	36 minutes
3	10: 20	12 minutes

- FCFS

- $(24+74+105+112)/4=78.75$ min

- SJF

- $(24+122+57+16)/4=54.75$ min

Example 2

- 5 processes arrived almost simultaneously

P1 P2 P3 P4 P5

- For each of the scheduling algorithms, determine the **scheduling order** and the **average turnaround time**
 - FCFS
 - RR (time slice = 1)
 - SJF
 - Priority Scheduling

Process	Running Time	Priority
P1	8	4
P2	4	3
P3	6	2
P4	2	1
P5	10	5 (highest)

Example 2

- Average turnaround time
- FCFS
 - 17.6
- RR (time slice = 1)
 - 21
- SJF
 - 14
- Priority Scheduling
 - 21.6

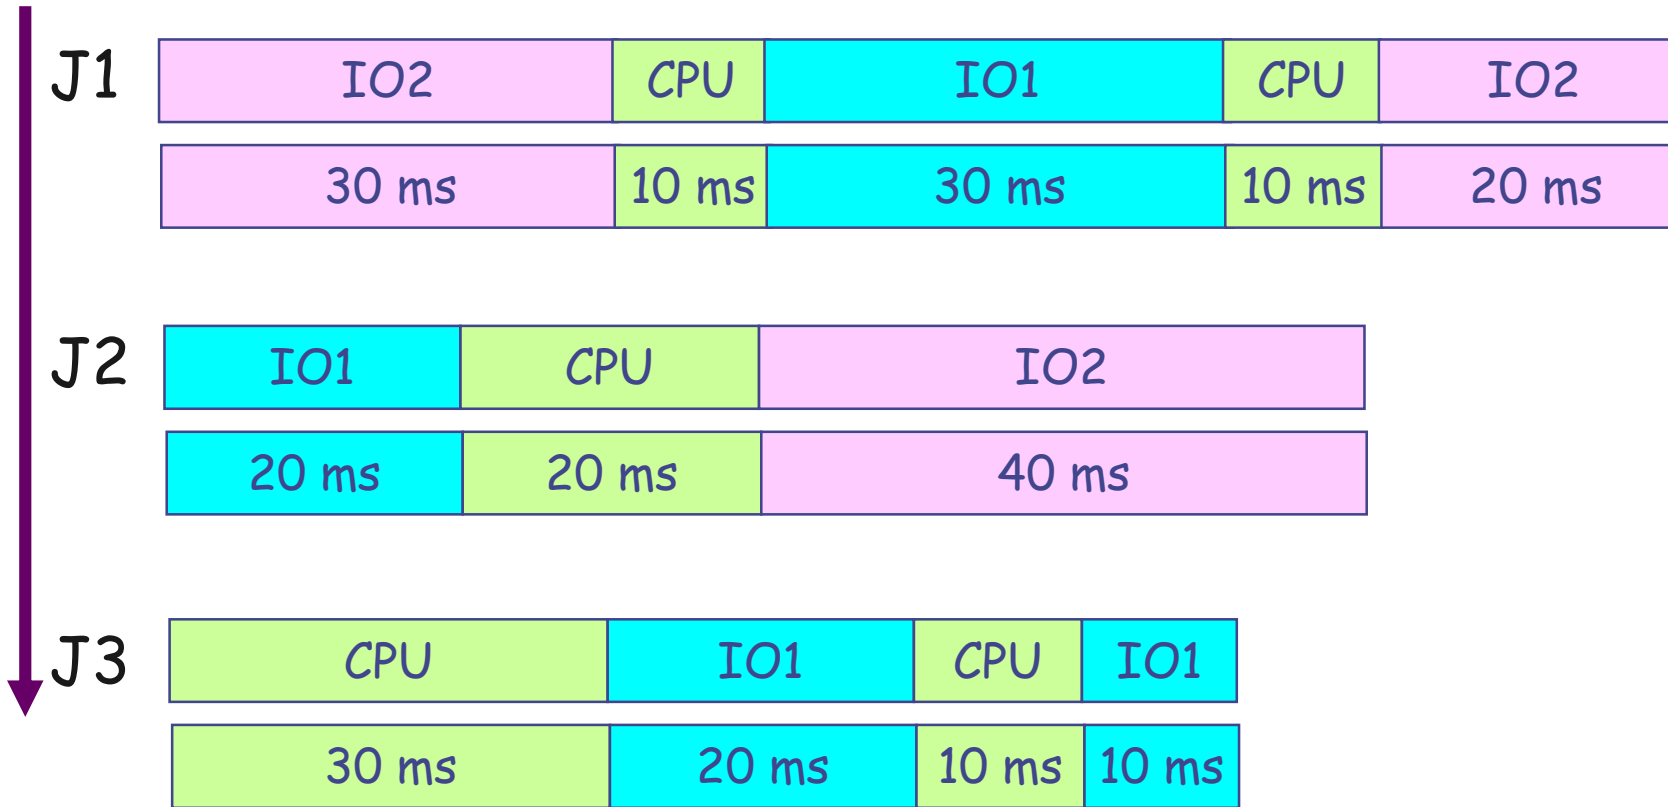
Process	Running Time	Priority
P1	8	4
P2	4	3
P3	6	2
P4	2	1
P5	10	5 (highest)

Example 3

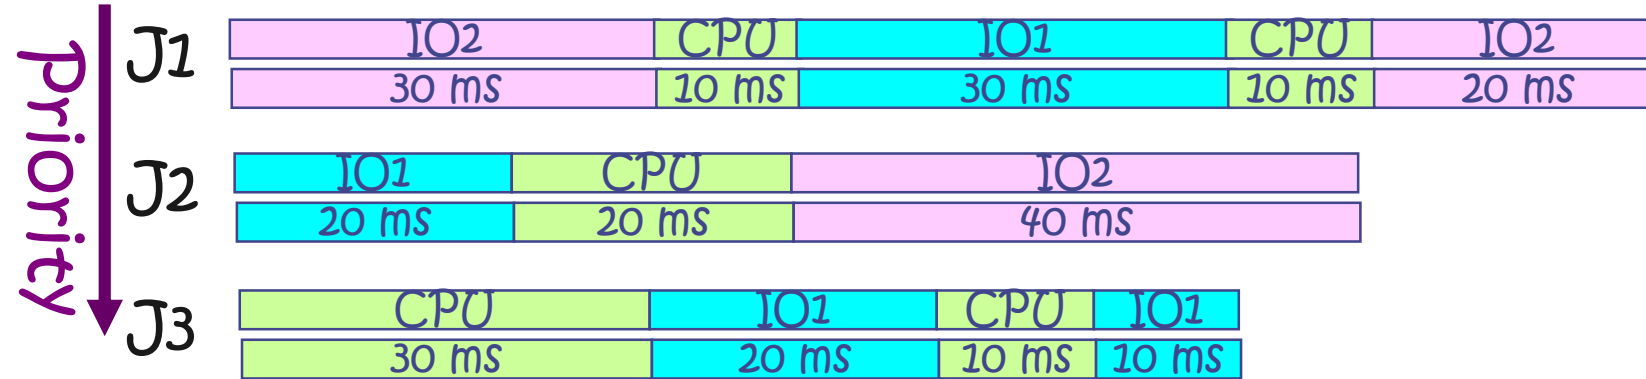
- In a multi programming system it has 1 processor and 2 IO devices. 3 jobs come at once, called J1、 J2 and J3. The running time queue they use as follows.
 - J1: IO2(30ms); CPU(10ms); IO1(30ms); CPU(10ms); IO2(20ms)
 - J2: IO1(20ms); CPU(20ms); IO2(40ms)
 - J3: CPU(30ms); IO1(20ms); CPU(10ms); IO1(10ms)
- Assume that CPU、 IO1 and IO2 work **parallel**, J1 has highest priority, J2 has second one, J3 has third. The higher priority job **can preempt the CPU**, but **can not preempt IO**.

Example 3

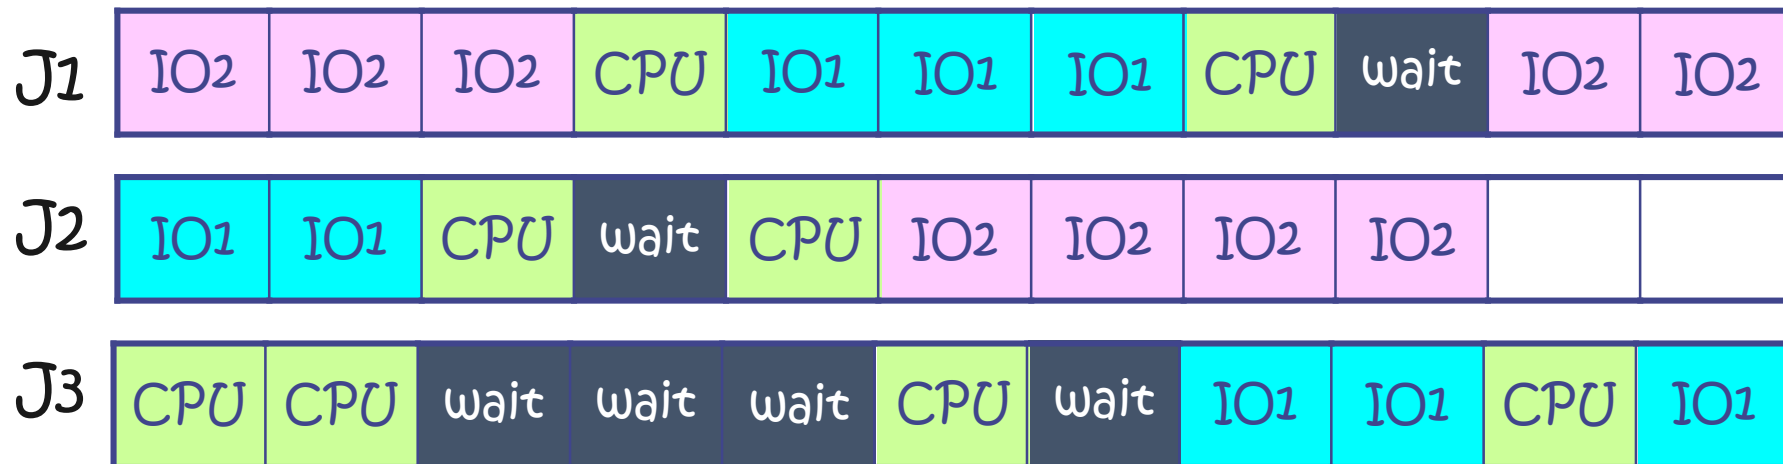
Priority



Example3



• 10ms

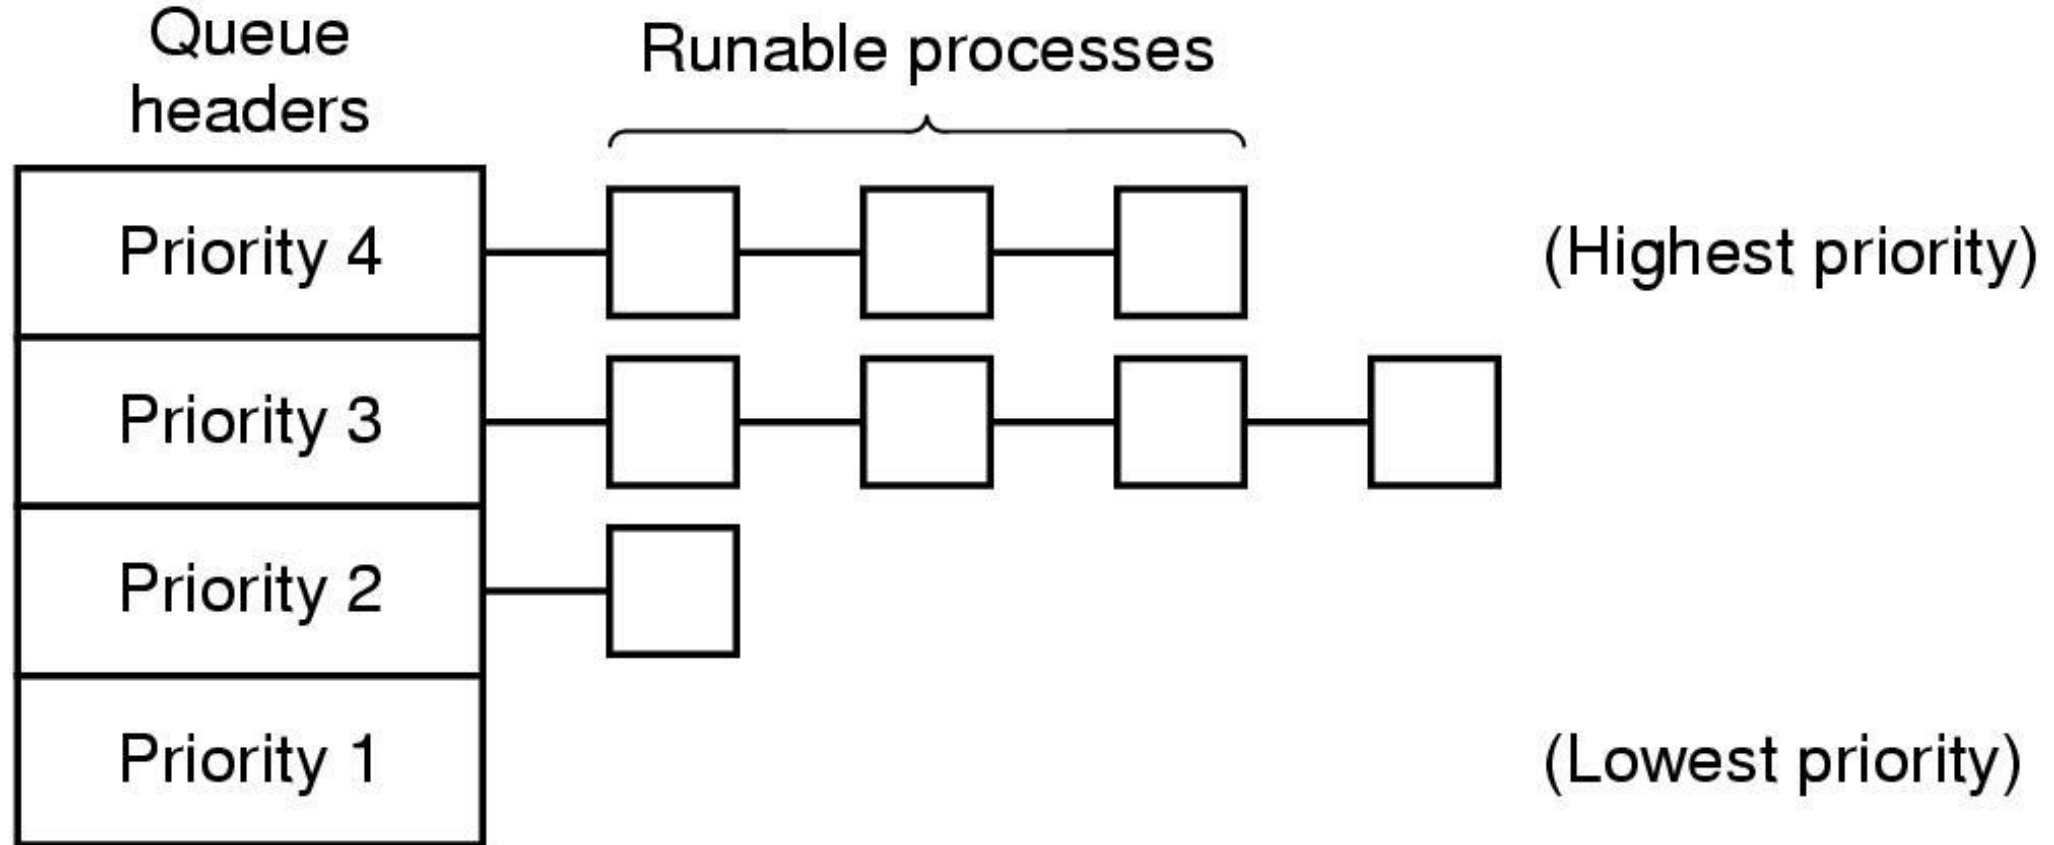


Example 3

10ms

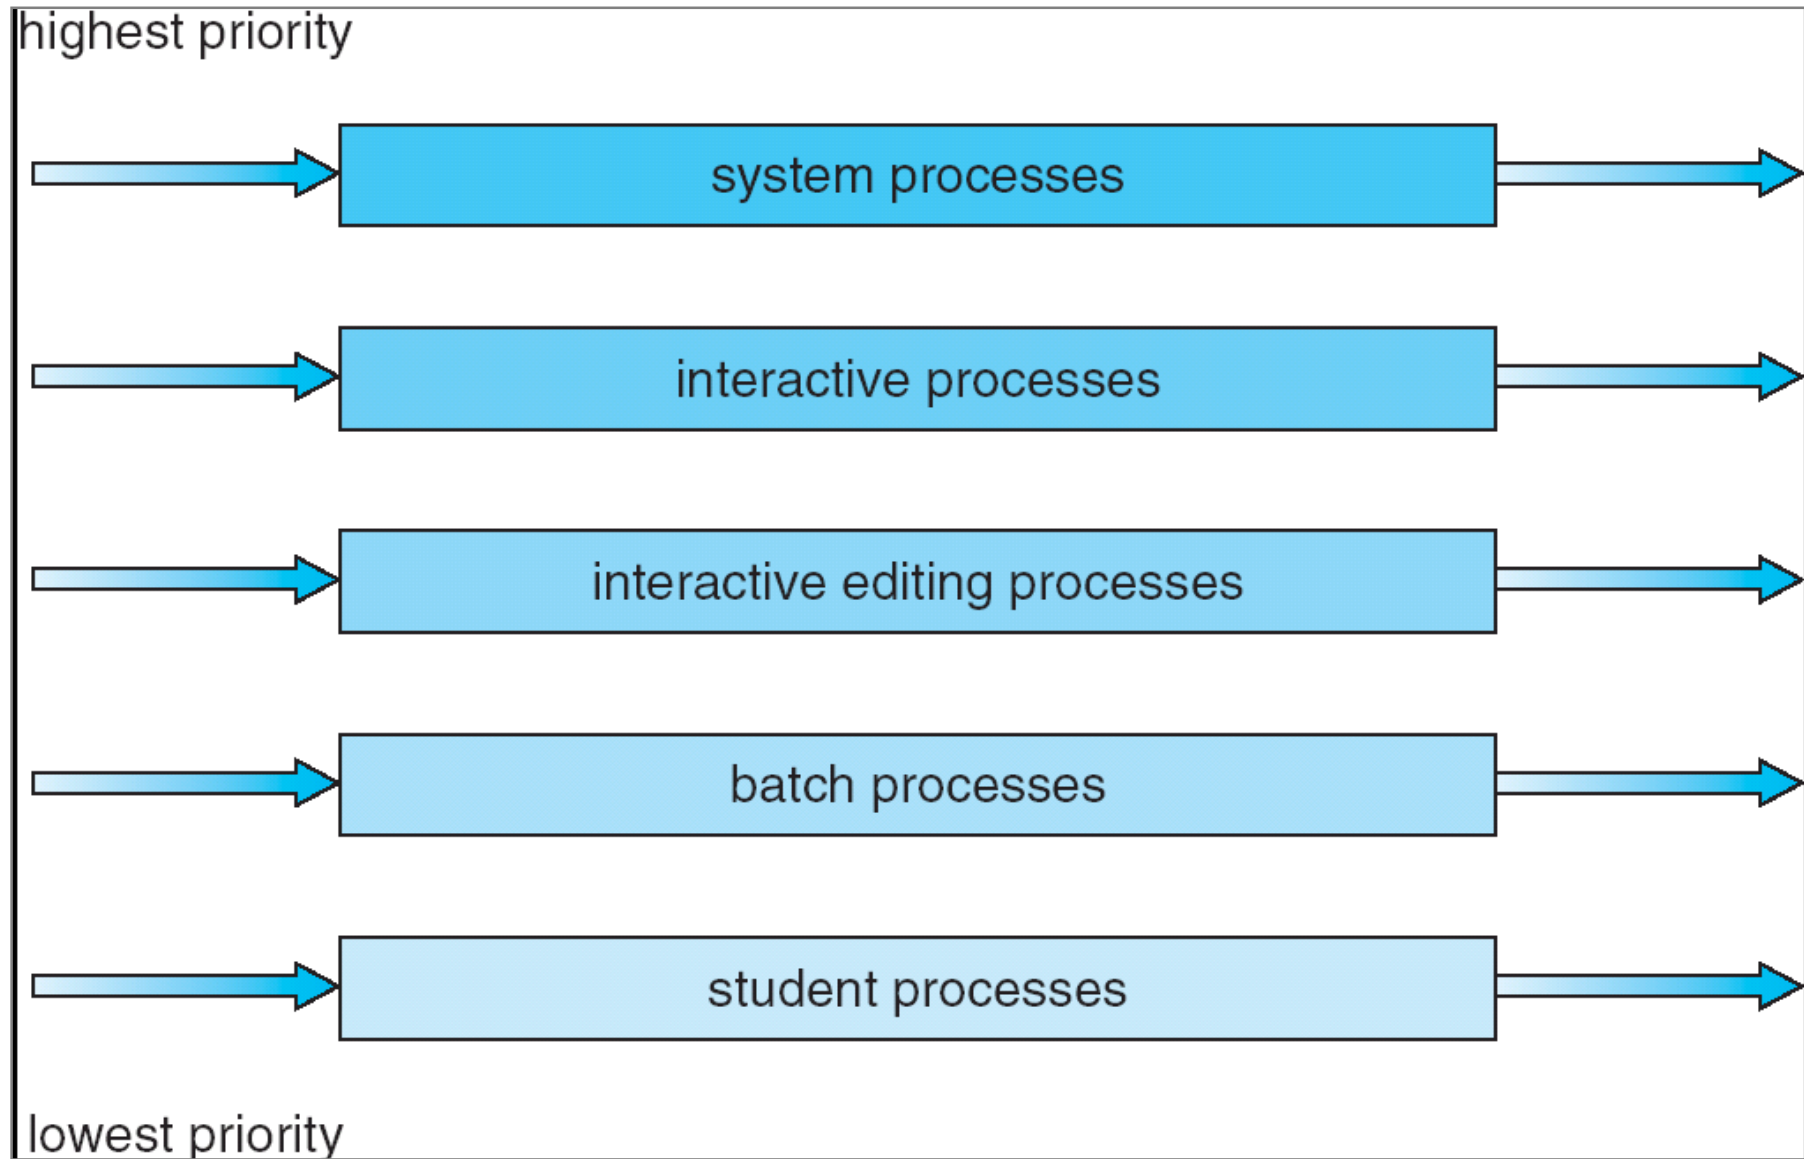


Multilevel Queue Scheduling





Multilevel Queue Scheduling



Multilevel Queue Scheduling

- Example:
 - Ready queue is partitioned into **separate queues**: foreground (interactive) / background (batch)
 - **Each queue has its own scheduling algorithm**
 - foreground – RR
 - background – FCFS

Multilevel Queue Scheduling

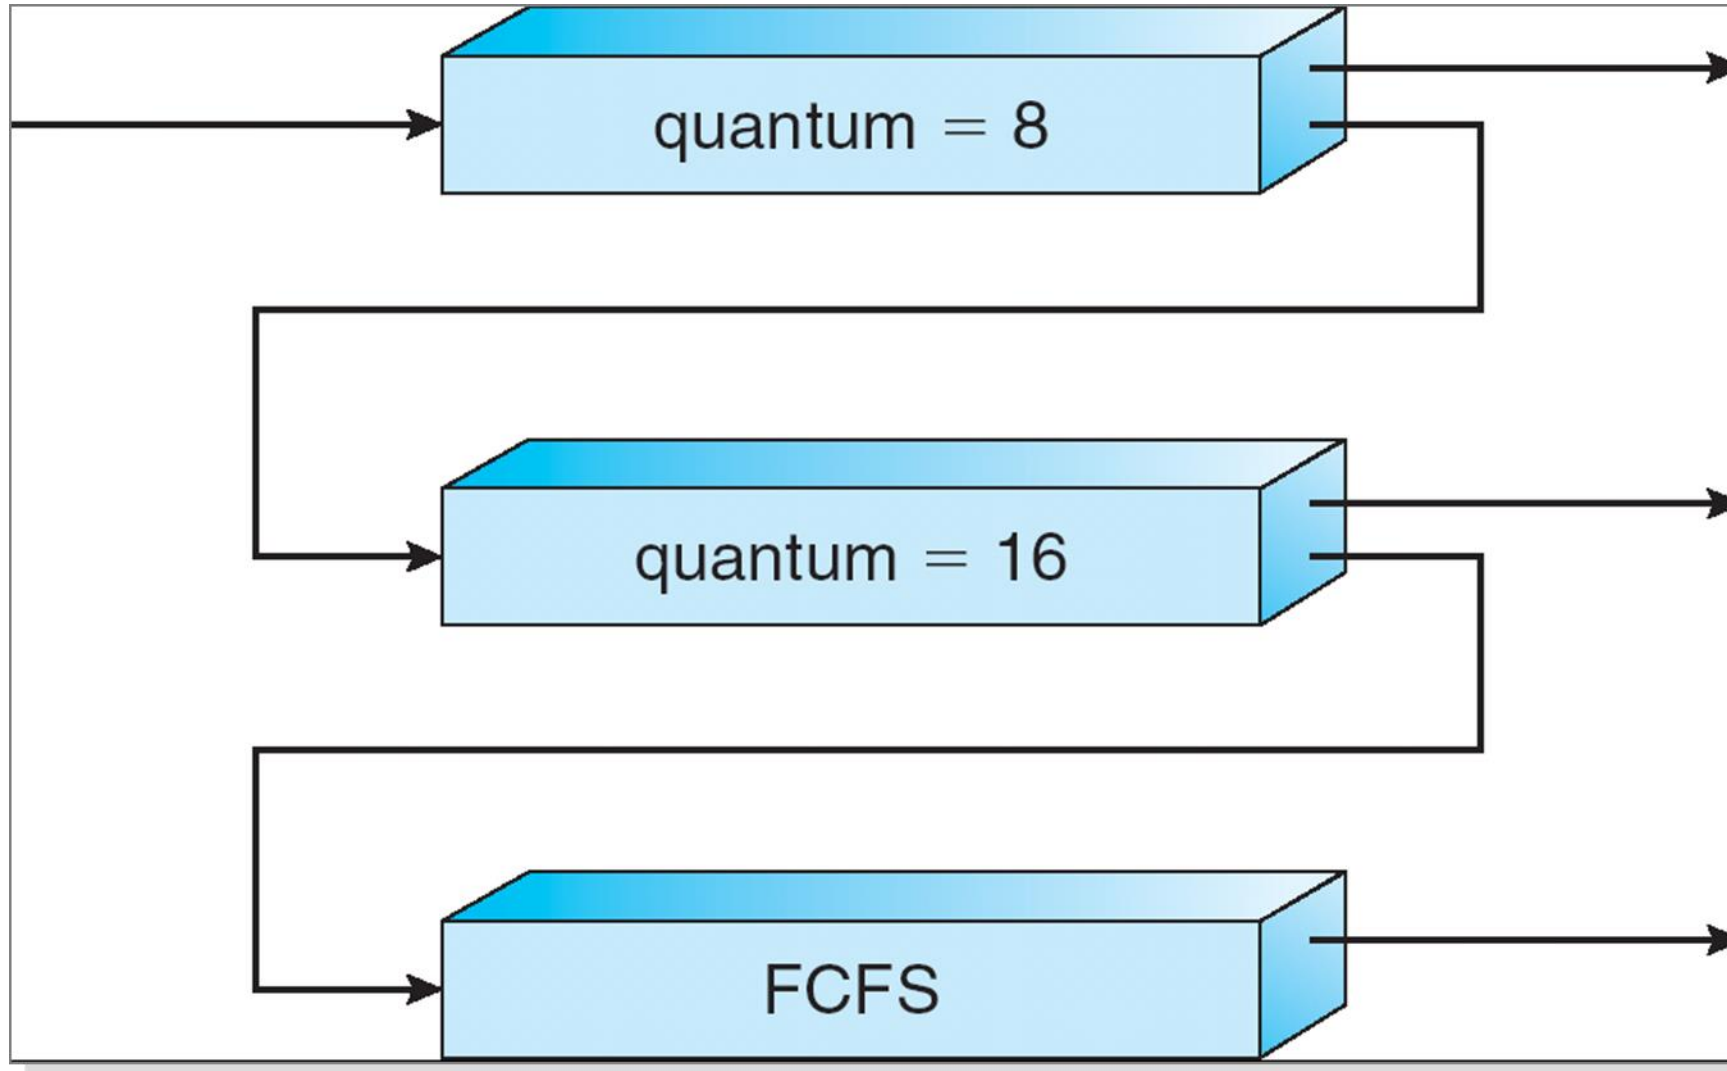
- Scheduling must be done between the queues
 - **Fixed priority scheduling**
 - serve all from highest priority, then next priority, etc.
 - e.g., serve all from foreground then from background.
 - Possibility of starvation.
 - **Time slice**
 - each queue gets a certain amount of CPU time which it can schedule amongst its processes
 - e.g., 80% to foreground in RR, 20% to background in FCFS

Multilevel Feedback Queue Scheduling



- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler is defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method to determine *when to upgrade a process*
 - method to determine *when to demote a process*
 - method to determine *which queue a process will enter when that process needs service*


Example of Multilevel Feedback Queue



Example of Multilevel Feedback Queue

- Three queues:

- Q_0 – RR with time quantum 8 milliseconds
- Q_1 – RR time quantum 16 milliseconds
- Q_2 – FCFS



- Scheduling

- A new job enters queue Q_0 which is served RR. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
- At Q_1 job is again served RR and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .

Multilevel Feedback Queue

- I/O bounded
 - The queue with highest priority
- CPU bounded
 - Lower and lower
- No need to estimate the length of the next CPU bursts

Scheduling Details

- **Countermeasure**: user action that can foil intent of the OS designer
 - To avoid moving to a lower priority queue, a user may put in a bunch of meaningless I/O , e.g., output a meaningless char on the screen regularly, to keep its priority high
 - Of course, if everyone did this, wouldn't work!



Other issues

Multiple-Processor Scheduling

- CPU scheduling more complex when multiple CPUs are available
- **Homogeneous processors** within a multiprocessor system
 - Only homogeneous systems are discussed here
- Scheduling
 1. Each processor has a separate queue
 2. All processes use a common ready queue
 - a. Each processor is self-scheduling
 - b. Appoint a processor as scheduler for the other processes (the master-slave structure)

Real-Time Scheduling

- Timesharing vs Realtime
- **Hard real-time** systems – required to complete a critical task within a guaranteed amount of time
- **Soft real-time** computing – requires that critical processes receive priority over less fortunate ones

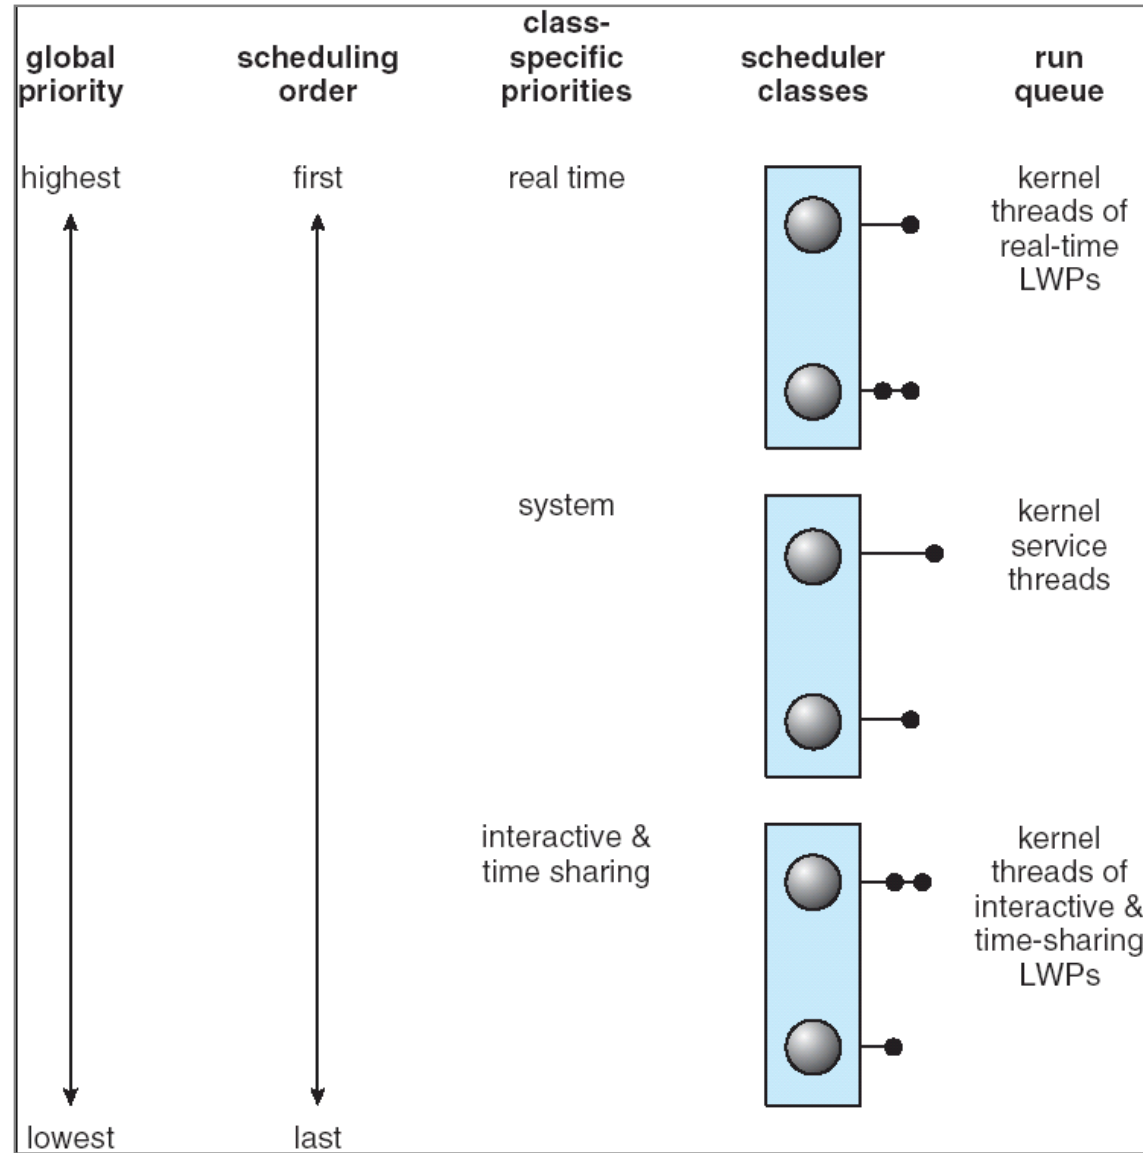


Operating Systems Examples

Operating System Examples

- Solaris scheduling
- Windows XP scheduling
- Linux scheduling

Solaris Scheduling (Priority-based)



Windows XP Priorities (priority-based, preemptive)



	real-time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1

Win32 APIs

- Suspend/ResumeThread: 挂起正在运行的线程或激活一个暂停运行的线程。
- Get/SetPriorityClass: 读取或设置进程的基本优先级类型
- Get/SetThreadPriority: 读取或设置线程相对优先级
- Get/SetProcessPriorityBoost: 读取或设置当前进程缺省优先级提升控制
- Get/SetThreadPriorityBoost: 读取或设置暂时提升线程优先级状态
- Get/SetProcessAffinityMask: 读取或设置进程的亲合处理器集合
- SetThreadAffinityMask: 设置线程的亲合处理器集合，只允许该线程在指定处理器集合运行
- SetThreadIdealProcessor: 设置特定线程的首选处理器
- SwitchToThread: 当前线程放弃一个或多个时间配额的运行
- Sleep: 使当前线程等待指定的一段时间（单位为毫秒）。0表示放弃该线程的剩余时间配额。
- SleepEx: 使当前线程进入等待状态，直到I/O处理完成

Linux Scheduling

- Two algorithms: time-sharing and real-time
- Time-sharing
 - Prioritized credit-based – process with most credits is scheduled next
 - Credit subtracted when timer interrupt occurs
 - When credit = 0, another process is chosen
 - When all processes have credit = 0, recrediting occurs
 - Based on factors including priority and history
- Real-time
 - Soft real-time



Algorithm Evaluation

Algorithm Evaluation

- Criteria to select a CPU scheduling algorithm may include several measures, such as:
 - **Maximize CPU utilization** under the constraint that the maximum response time is 1 second
 - **Maximize throughput** such that turnaround time is (on average) linearly proportional to total execution time

How to Evaluate a Scheduling Algorithm?



- Deterministic modeling
- Queuing models
- Simulation/Implementation

Algorithm Evaluation

- **Deterministic modeling**
 - Analytic evaluation
 - Input: a given **algorithm** and a **system workload**
 - Output: performance of the algorithm for that workload

Algorithm Evaluation

- **Deterministic modeling**
 - A simple and fast method. It gives the exact numbers, allows the algorithms to be compared.
 - It requires exact numbers of input, and its answers apply to only those cases. In general, deterministic modeling is too specific, and requires too much exact knowledge, to be useful.
 - Usage
 - Describing algorithm and providing examples
 - A set of programs that may run over and over again
 - Indicating the trends that can then be proved

Algorithm Evaluation

- **Queuing Models** - Queuing network analysis
 - Using
 - the **distribution** of service times (CPU and I/O bursts)
 - the **distribution** of process arrival times
 - The computer system is described as a network of servers. Each server has a queue of waiting processes.
 - Determining
 - **utilization, average queue length, average waiting time, etc.**

Algorithm Evaluation

• Simulations

- Simulations involve programming a model of the system. Software data structures represent the major components of the system.
- Coding a simulator can be a major task
- Generating data to drive the simulator
 - a random number generator.
 - **trace tapes**: created by monitoring the real system
- **More accurate**
- Expensive (several hours of computer time).
- Large storage
 - Recording the sequence of actual events.

Algorithm Evaluation

- **Implementation**

- Put data into a real system and see how it works.
- The **only accurate** way
 1. cost is too high
 2. environment will change (All methods have this problem!)
 - e.g., To avoid moving to a lower priority queue, a user may output a meaningless char on the screen regularly to keep itself in the interactive queue.



Summary

Summary

- **Scheduling**: selecting a waiting process from the ready queue and allocating the CPU to it
- **FCFS Scheduling**:
 - Run processes to completion in order of submission
 - Pros: Simple
 - Cons: Short jobs get stuck behind long ones
- **Round-Robin Scheduling**:
 - Give each process a small amount of CPU time when it executes; cycle between all ready processes
 - Pros: Better for short jobs
 - Cons: Poor when jobs are same length

Summary

- **Shortest Job First (SJF)/Shortest Remaining Time First (SRTF):**
 - Run whatever job has the least amount of computation to do/least remaining amount of computation to do
 - Pros: optimal (average response time)
 - Cons: hard to predict future, unfair
- **Multi-Level Feedback Queue Scheduling:**
 - Multiple queues of different priorities
 - Automatic promotion/demotion of process priority in order to approximate SJF/SRTF

Summary

- **Countermeasure**: user action that can foil intent of the OS designer
- **Evaluation of mechanisms**:
 - Analytical/Deterministic modeling
 - Queuing Theory
 - Simulation
 - Implementation

Process Schedule

