



Process Synchronization

Outline

- Background
- The **Critical-Section** Problem
- Synchronization Hardware
- **Semaphores**
- Classic Problems of Synchronization
- Monitors
- Synchronization Examples

"Too much milk"

- People need to coordinate:

Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away



Motivation: “Too much milk”

- Great thing about OS's – analogy between problems in OS and problems in real life
 - Help you understand real life problems better
 - But, computers are **much stupider** than people





Background

Background

- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

Race Condition

- Producer and Consumer
- `count++` could be implemented as
 - `register1 = count`
 - `register1 = register1 + 1`
 - `count = register1`
- `count--` could be implemented as
 - `register2 = count`
 - `register2 = register2 - 1`
 - `count = register2`

Race Condition

- Assume count is initially 5.

count++	count--
register1 = count	register2 = count
register1 = register1 + 1	register2 = register2 - 1
count = register1	count = register2

Consider this execution interleaving with “count = 5” initially:

S0: producer	execute	register1 = count	{register1 = 5}
S1: producer	execute	register1 = register1 + 1	{register1 = 6}
S2: consumer	execute	register2 = count	{register2 = 5}
S3: consumer	execute	register2 = register2 - 1	{register2 = 4}
S4: producer	execute	count = register1	{count = 6 }
S5: consumer	execute	count = register2	{count = 4}

Race Condition

- We're off.
- P gets off to an early start
- C says “humph, better go fast” and tries really hard
- P goes ahead and writes “6”
- C goes and writes “4”
- P says “HUH??? I could have sworn I put a 6 there”

Consider this execution interleaving with “count = 5” initially:

S0: producer	execute	register1 = count	{register1 = 5}
S1: producer	execute	register1 = register1 + 1	{register1 = 6}
S2: consumer	execute	register2 = count	{register2 = 5}
S3: consumer	execute	register2 = register2 - 1	{register2 = 4}
S4: producer	execute	count = register1	{count = 6 }
S5: consumer	execute	count = register2	{count = 4}

Race Condition

- We're off.
- P gets off to an early start
- C says “humph, better go fast” and tries really hard
- P goes ahead and writes “6”
- C goes and writes “4”
- P says “HUH??? I could have sworn I put a 6 there”

Could this happen on a uniprocessor?

Yes! Unlikely, but if you depending on it not happening, it will and your system will break...

Definitions

- **Race condition:** The situation where several processes access and manipulate shared data concurrently. The final value of the shared data depends upon which process finishes last.
- To prevent race conditions, concurrent processes must be **synchronized**.

Definitions

- **Synchronization:**

Using **atomic operations** to ensure cooperation between processes

For now, only loads and stores are **atomic**

We are going to show that it's hard to build anything useful with **only** loads and stores

Definitions

- **Atomic Operation**: an operation that always runs to completion or not at all
 - It is **indivisible**: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
 - **Fundamental** building block – if no atomic operations, then have no way for processes/threads to work together
- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic

Definitions

- **Mutual Exclusion**: ensuring that only one process does a particular thing at a time
 - One process **excludes** the other while doing its task
- **Critical Section**: piece of code that only one process can execute at once. Only one process at a time will get into this section of code.
 - Critical section is the result of mutual exclusion
 - Critical section and mutual exclusion are two ways of describing the same thing.



The Critical-Section Problem

The Critical-Section Problem

- n processes are all competing to use some **shared data**
- Each process has a **code segment**, called **critical section**, in which the shared data is accessed.
- **Problem** – ensure that when one process is executing in its critical section, **no other process is allowed to execute in its critical section.**

The Critical-Section Problem

```
do {  
    critical section  
  
    remainder  
    section  
  
} while (TRUE);
```

```
do {  
    buy milk;  
  
    be in a daze;  
  
} while (TRUE);
```

Too Much Milk: Correctness Properties


- Need to be careful about correctness of concurrent programs, since non-deterministic
 - Always write down behavior first
 - Impulse is to start coding first, then when it doesn't work, pull hair out
 - Instead, think first, then code
- What are the correctness properties for the “Too much milk” problem???
- Never more than one person buys
- Someone buys if needed
- Restrict ourselves to use only atomic load and store operations as building blocks

Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of “lock”)
 - Remove note after buying (kind of “unlock”)
 - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (no Milk) {  
    if (no Note) {  
        leave Note;  
        buy milk;  
        remove Note;  
    }  
}
```

- Result?




Too Much Milk: Solution #1

Process A	Process B
if (no Milk) if (no Note)	
	if (no Milk) if (no Note)
leave Note; buy milk; remove Note;	
	leave Note; buy milk; remove Note;

Too Much Milk: Solution #1

- Result?
 - Still too much milk but only **occasionally**!
 - Process can get context switched after checking milk and note but before buying milk!
- Solution makes problem worse since fails **intermittently**
 - Makes it really hard to debug...



Too Much Milk: Solution #1 $\frac{1}{2}$

- Clearly the Note is not quite blocking enough
 - Let's try to fix this by placing note first
- Another try at previous solution:

```
leave Note;  
if (no Milk) {  
    if (no Note) {  
        buy milk;  
    }  
}  
remove note;
```

```
if (no Milk) {  
    if (no Note) {  
        leave Note;  
        buy milk;  
        remove Note;  
    }  
}
```

- What happens here?
 - Well, with human, probably nothing bad
 - With computer: no one ever buys milk



Too Much Milk Solution #2

- How about labeled notes?
 - Now we can leave note before checking
- Algorithm looks like this:

Process A

```
leave note A;  
  if (no Note B) {  
    if (no Milk) {  
      buy Milk;  
    }  
  }  
remove note A;
```

Process B

```
leave note B;  
  if (no Note A) {  
    if (no Milk) {  
      buy Milk;  
    }  
  }  
remove note B;
```

- Does this work?

Too Much Milk Solution #2


Process A	Process B
leave note A;	
	leave note B;
if (no Note B) if (no Milk) buy Milk;	
	if (no Note A) if (no Milk) buy Milk;
remove note A;	
	remove note B;

Too Much Milk Solution #2

- Possible for neither process to buy milk
 - Context switches at exactly the wrong times can lead each to think that the other is going to buy
- Really insidious:
 - **Extremely unlikely** that this would happen, but will at worse possible time

Too Much Milk Solution #2: problem!

- I'm not getting milk, You're getting milk
- This kind of lockup is called “starvation!”



Too Much Milk Solution #3

- Here is a possible two-note solution:

```
Process A  
leave note A;  
while (note B) {  
    do nothing;  
}  
if (no Milk) {  
    buy milk;  
}  
remove note A;
```

```
Process B  
leave note B;  
    if (no Note A) {  
        if (no Milk) {  
            buy milk;  
        }  
    }  
remove note B;
```

- Does this work?

Too Much Milk Solution #3

- Here is a possible two-note solution:

Process A

```
leave note A;  
  while (note B) //X  
    do nothing;  
if (no Milk)  
  buy milk;  
remove note A;
```

Process B

```
leave note B;  
if (no Note A) //Y  
  if (no Milk)  
    buy milk;  
remove note B;
```

- Does this work? Yes. Both can guarantee that:
 - It is safe to buy, or other will buy, ok to quit
- At X:
 - if no note B, safe for A to buy, otherwise wait to find out what will happen
- At Y:
 - if no note A, safe for B to buy, otherwise, A is either buying or waiting for B to quit

Solution #3 discussion

- Our solution protects a single “Critical-Section” piece of code for each process:

```
if (no Milk) {  
    buy milk;  
}
```

- Solution #3 works, but it's really unsatisfactory
 - Really complex – even for this simple an example
 - Hard to convince yourself that this really works
 - A's code is different from B's – what if lots of processes?
 - Code would have to be slightly different for each process
 - While A is waiting, it is consuming CPU time
 - This is called “busy-waiting”

Solution #3 discussion

- There's a better way
 - Have hardware provide better (higher-level) primitives than atomic load and store
 - Build even higher-level programming abstractions on this new hardware support

Too Much Milk: Solution #4

- **Lock**: prevents someone from doing something
- Suppose we have some sort of implementation of a lock.
 - `Lock.Acquire()` – Wait until lock is free, then grab
 - `Lock.Release()` – Unlock, waking up anyone waiting
 - These must be atomic operations – if two processes are waiting for the lock and both see it's free, only one succeeds to grab the lock

Too Much Milk: Solution #4

- For example: fix the milk problem by putting a key on the refrigerator
 - Lock it and take key if you are going to go buy milk
 - Fixes too much: person B angry if only wants other juice



- Of course – We don't know how to make a lock yet

Too Much Milk: Solution #4

- Then, our milk problem is easy:


```
milklock.Acquire();  
if (no milk)  
    buy milk;  
milklock.Release();
```

- Once again, section of code between `Acquire()` and `Release()` called a “**Critical Section**”

Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (TRUE);
```

Solution to Critical-section Problem Using Locks



How to implement Locks?



- **Lock**: prevents someone from doing something
 - Lock before entering critical section and before accessing shared data
 - Unlock when leaving, after accessing shared data
 - Wait if locked
- Atomic Load/Store: get solution like Milk #3
 - Pretty complex and error prone



Synchronization Hardware

Naive use of Interrupt Enable/Disable



- Naive Implementation of locks:

```
LockAcquire { disable Ints; }  
LockRelease { enable Ints; }
```

- Problems with this approach:

- **Can't let user do this!** Consider following:

```
LockAcquire();  
While(TRUE) { ; }
```

- What happens with I/O or other important events?
 - “Reactor about to meltdown. Help?”
- Generally too inefficient on multiprocessor systems
 - Disabling interrupts on all processors requires messages and would be very time consuming
 - Operating systems using this not broadly scalable

Atomic Read-Modify-Write instructions

- Alternative: atomic instruction sequences
 - These instructions read a value from memory and write a new value atomically
 - Unlike disabling interrupts, can be used on both uniprocessors and multiprocessors
 - Many systems provide hardware support for critical section code

TestAndSet Instruction

- Definition:

```
boolean TestAndSet (boolean *target) {  
    boolean rv = *target;  
    *target = TRUE;  
    return rv;  
}
```


Implementing Locks with TestAndSet

- A flawed, but simple solution:

```
int value = 0; // Free
Acquire() {
    while (TestAndSet(&value)); // while busy
}
Release() {
    value = 0;
}
```

- Simple explanation:
 - If lock is free, *TestAndSet* reads 0 and sets value=1, so lock is now busy. It returns 0 so while exits.
 - If lock is busy, *TestAndSet* reads 1 and sets value=1 (no change). It returns 1, so while loop continues
 - When we set value = 0, someone else can get lock
- **Busy-Waiting**: process consumes cycles while waiting

Problem: Busy-Waiting for Lock

- Positives for this solution
 - Machine can receive interrupts
 - User code can use this lock
 - Works on a multiprocessor
- Negatives
 - This is very inefficient because the **busy-waiting** process will consume cycles waiting
 - **Priority Inversion**: If busy-waiting process has higher priority than process holding lock \Rightarrow no progress!



Swap Instruction

- Definition:

```
void Swap (boolean *a, boolean *b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```


Solution using Swap

- Shared Boolean variable `lock` initialized to FALSE
- Each process has a local Boolean variable `key`
- Solution:

```
do {  
    key = TRUE;  
    while (key == TRUE)  
        Swap (&lock, &key);  
  
    //critical section  
  
    lock = FALSE;  
  
    //remainder section  
  
} while (TRUE);
```




Semaphores



Semaphores

- Semaphore from railway analogy
 - Here is a semaphore initialized to 2 for resource control:



Semaphores

- Semaphores are a kind of generalized lock
 - First defined by Dijkstra in 1965
 - Main synchronization primitive used in original UNIX
 - Does not require busy waiting
 - (if be implemented using wait queue)
 - Less complicated



Semaphores

- Definition: a Semaphore has **a non-negative integer value** and supports the following two operations (spin lock):
 - **P()**: an **atomic** operation that **waits** for semaphore to become positive, then decrements it by 1
 - Think of this as the **wait()** operation
 - **V()**: an **atomic** operation that increments the semaphore by 1, waking up a waiting P, if any
 - Think of this as the **signal()** operation



Semaphores


- The definition of **wait ()** is as follows:

```
wait (S) {  
    while (S <= 0)  
        ; // no-op  
    S--;  
}
```

- The definition of **signal ()** is as follows:

```
signal (S) {  
    S++;  
}
```





Semaphore

Semaphores

- Note that **P** stands for “**proberen**” (to test) and **V** stands for “**verhogen**” (to increment) in Dutch
 - $P(S)$ and $V(S)$
 - $\text{Wait}(S)$ and $\text{Signal}(S)$
 - $S.\text{wait}()$ and $S.\text{signal}()$



Semaphores

- No negative values
- **Only** operations allowed are P and V – can't read or write value, except to set it initially
- Operations must be **atomic**
 - Two P's together can't decrement value below zero
 - Similarly, process going to sleep in P won't miss wakeup from V – even if they both happen at same time

Semaphore Implementation

- The main disadvantage of the semaphore definition given here is that it requires **busy waiting**.
 - While a process is in its critical section, any other process that tries to enter its critical section must loop continuously in the entry code.
- This type of semaphore is also called a **spin lock**.

Semaphores

- The definition of **wait** () is as follows:



```
wait (S) {  
    while (S <= 0)  
        ; // no-op  
    S--;  
}
```

busy
waiting

- The definition of **signal** () is as follows:

```
signal (S) {  
    S++;  
}
```







Semaphore Implementation with no Busy Waiting



- With each semaphore there is an associated waiting queue.



- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue.
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.



Semaphore Implementation with no Busy Waiting

- Implementation of wait:

```
Wait (S){  
    value--;  
    if (value < 0) {  
        add this process to waiting queue  
        block();  
    }  
}
```

- Implementation of signal:

```
Signal (S){  
    value++;  
    if (value <= 0) {  
        remove a process P from the waiting queue  
        wakeup(P);  
    }  
}
```

Two Types of Semaphores

- **Binary** semaphore – integer value can range only between 0 and 1; can be simpler to implement.
 - Also known as **mutex locks**
- **Counting** semaphore – integer value can range over an unrestricted domain.
 - Can implement a counting semaphore **S** as a binary semaphore.

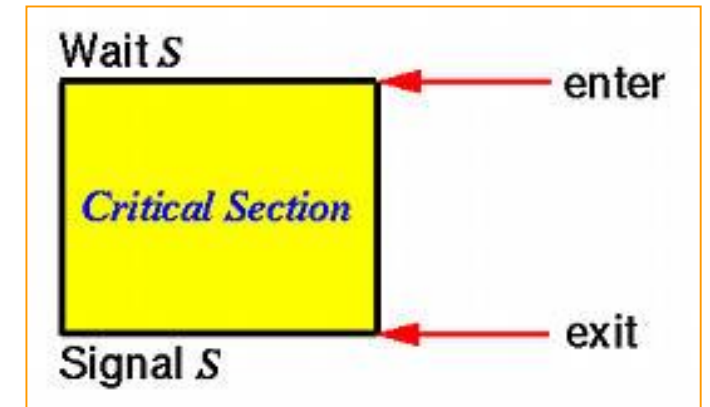
Two Types of Semaphores

- Provides mutual exclusion


```
Semaphore mutex; // initialized to 1  
do {
```

```
    wait (mutex);  
    // Critical Section  
    signal (mutex);  
    // remainder section
```

```
} while (TRUE);
```




Semaphore




Precedence Graph

- A DAG (Directed Acyclic Graph)




Semaphore

- Semaphore a, b, c, d, e, f, g;
- $a=0, b=0, c=0, d=0, e=0, f=0, g=0$;
- $S1() \{ \dots; V(a); V(b); \}$
- $S2() \{ P(a); \dots; V(c); V(d); \}$
- $S3() \{ P(b); \dots; V(e); \}$
- $S4() \{ P(c); \dots; V(f); \}$
- $S5() \{ P(d); \dots; V(g); \}$
- $S6() \{ P(f); P(g); P(e); \dots; \}$



Semaphore

```
While (1) {  
    P (SA) ;  
    ping;  
    V (SB) ;  
}
```



```
While (1) {  
    P (SB) ;  
    pong;  
    V (SA) ;  
}
```

Deadlock and Starvation

- **Deadlock** – two or more processes are waiting **indefinitely** for an event that can be caused by only one of the waiting processes
- Let **S** and **Q** be two semaphores initialized to 1

P_0
wait (S);
wait (Q);
...
signal (S);
signal (Q);

P_1
wait (Q);
wait (S);
...
signal (Q);
signal (S);

Deadlock and Starvation

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.