



# Classic Problems of Synchronization


# Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem
- The three problems are important, because they are
  - examples for a large class of concurrency-control problems.
  - used for testing nearly every newly proposed synchronization scheme.
- **Semaphores** are used for synchronization in our solutions.

# Producer-consumer problem with a bounded buffer

- **Problem Definition**

- Producer puts things into a shared buffer
- Consumer takes them out
- Need synchronization to coordinate producer/consumer



- **Put a fixed-size buffer between them**

- Need to synchronize access to this buffer
- Producer needs to wait if buffer is full
- Consumer needs to wait if buffer is empty

- **Example: Coke machine**

- Producer can put limited number of cokes in machine
- Consumer can't take cokes out if machine is empty



# Bounded-Buffer

- Suppose that we wanted to provide a solution to the consumer-producer problem with a **bounded-buffer**.
- We can do so by having an integer **count** that keeps track of the number of products in the buffer.
- Initially, **count** is set to 0.
  - It is incremented by the producer after it produces a new product.
  - It is decremented by the consumer after it consumes a product.

# Bounded-Buffer

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int count = 0;
```

# Producer and Consumer

```
while (true) {  
    // produce an item and put in nextProduced  
    while (count == BUFFER_SIZE)  
        ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

```
while (true) {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
    //consume the item in nextConsumed  
}
```

# Producer and Consumer

```
while (true) {  
    while (count == BUFFER_SIZE)  
        ; // do nothing  
    buffer[in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```



```
while (true) {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
}
```



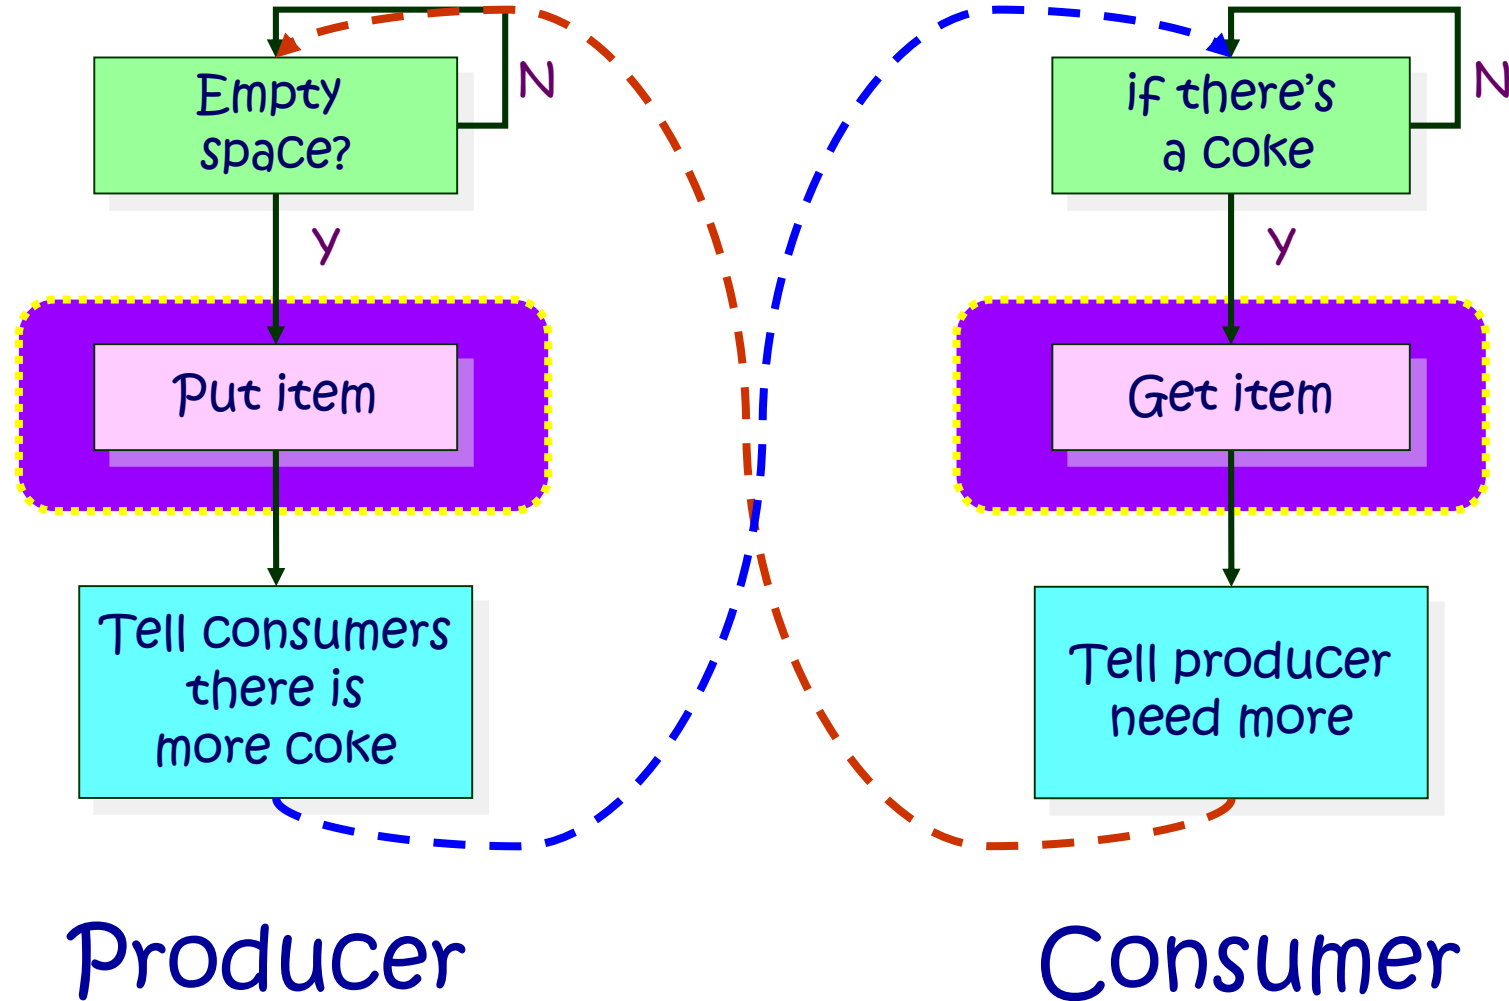
# Correctness constraints for solution

- **Correctness Constraints:**

- Consumer must wait for producer to fill the buffer, if no product in the buffer (scheduling constraint)
- Producer must wait for consumer to empty the buffer, if the buffer is full (scheduling constraint)
- Only one process can manipulate the buffer at a time (mutual exclusion)



# Correctness constraints for solution



# Correctness constraints for solution

- General rule of thumb: Use a separate semaphore for each constraint
  - Semaphore full;  
// consumer's constraint  
// initialized to the value 0
  - Semaphore empty;  
// producer's constraint  
// initialized to the value N.
  - Semaphore mutex;  
// mutual exclusion  
// initialized to the value 1

# Full Solution

```
Semaphore full = 0;           // Initially, no coke
Semaphore empty = N;          // Initially, num empty slots
Semaphore mutex = 1;           // No one using machine


Producer(item) {
    P(empty);                  // Wait until space
    P(mutex);                  // Wait until machine free
    Enqueue(item);
    V(mutex);
    V(full);                   // Tell consumers there is more coke
}

Consumer() {
    P(full);                   // Check if there's a coke
    P(mutex);                  // Wait until machine free
    item = Dequeue();
    V(mutex);
    V(empty);                  // Tell producer need more
    return item;
}
```

# Full Solution

- The structure of the **producer** process

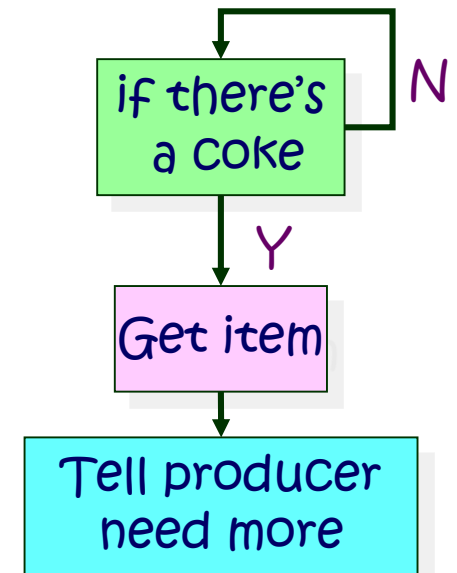
```
while (true) {  
    //    produce an item  
  
    wait (empty);  
    wait (mutex);  
  
    //    add the item to the buffer  
  
    signal (mutex);  
    signal (full);  
}
```



# Full Solution

- The structure of the **consumer** process

```
while (true) {  
  
    wait (full);  
    wait (mutex);  
  
    //  remove an item from buffer  
  
    signal (mutex);  
    signal (empty);  
  
    //  consume the removed item  
}
```



# Discussion about Solution

- Why asymmetry?
  - Producer does:  $P(\text{empty})$  ,  $V(\text{full})$
  - Consumer does:  $P(\text{full})$  ,  $V(\text{empty})$
- Is the order of P's important?
- Is the order of V's important?

```
Producer(item) {  
    P(empty);  
    P(mutex);  
    Enqueue(item);  
    V(mutex);  
    V(full);  
}
```

```
Consumer() {  
    P(full);  
    P(mutex);  
    item = Dequeue();  
    V(mutex);  
    V(empty);  
    return item;  
}
```

# Discussion about Solution

Is the order of P's important?

- Yes! Can cause deadlock
- Why?

```
Semaphore full = 0;           // Initially, no coke
Semaphore empty = N;          // Initially, num empty slots
Semaphore mutex = 1;          // No one using machine
```

```
Producer(item) {
    P(mutex);           // Wait until buffer free
    P(empty);            // Wait until space
    Enqueue(item);
    V(mutex);
    V(full);             // Tell consumers there is
                        // more coke
}
```

```
Consumer() {
    P(full);             // Check if there's a coke
    P(mutex);            // Wait until machine free
    item = Dequeue();
    V(mutex);
    V(empty);            // tell producer need more
    return item;
}
```

# Discussion about Solution


---

- **Is order of V's important?**
  - No, except that it might affect scheduling efficiency
- What if we have 2 producers or 2 consumers?
  - Do we need to change anything?




# Readers/Writers Problem

- Motivation: Consider a shared database
  - Two classes of users:
    - Readers – only read the data set; they do **not** perform any updates.
    - Writers – can both read and write.
  - Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.




# Basic Readers/Writers Solution

- Correctness Constraints:
  - Readers can access database when no writers
  - Writers can access database when no readers or writers
  - Only one process manipulates state variables at a time
- Basic structure of a solution:
  - **Reader()**
    - Wait until no writers
    - Access database
    - Check out – wake up a waiting writer
  - **Writer()**
    - Wait until no active readers or writers
    - Access database
    - Check out – wake up waiting readers or writers



# Readers/Writers Problem

- Is using a single lock on the whole database sufficient?



# The first readers-writers problem

---

- Suppose we have a shared memory area with the constraints detailed above.
- It is possible to protect the shared data behind a **mutex**, in which case no process/thread can access the data at the same time.

# The first readers-writers problem

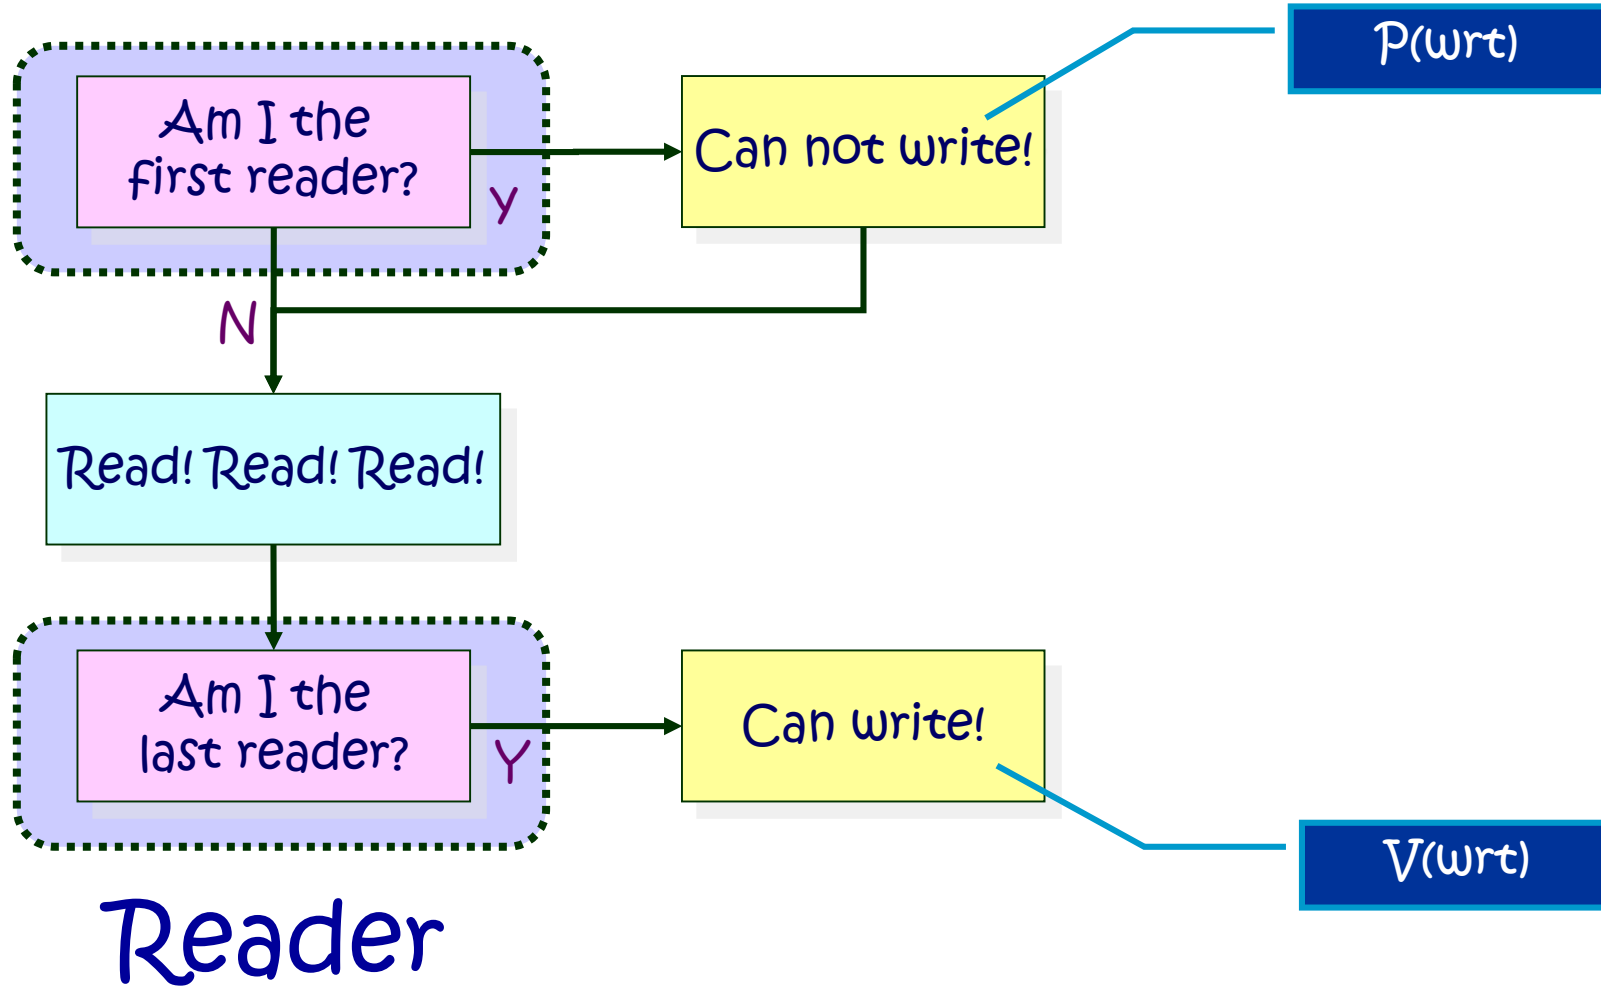
- The solution is sub-optimal.
  - Because it is possible that a reader R1 might have the lock, and then another reader R2 request access.
  - It would be foolish for R2 to wait until R1 was done before starting its own read operation.
  - Instead, R2 should start right away.
- This is the motivation for the **first readers-writers problem**, in which the constraint is added that **no reader shall be kept waiting if the share is currently opened for reading**.
- This is also called **readers-preference**.

# Readers-Writers Problem: Solution #1

- `int readcount = 0;`
- `semaphore mutex = 1, wrt = 1;`
- The structure of a writer process

```
while (true) {  
    wait (wrt);  
    // writing is performed  
    signal (wrt);  
}
```


# Readers-Writers Problem: Solution #1



# Readers-Writers Problem: Solution #1


- The structure of a reader process

```
while (true) {  
    wait(mutex);  
    readcount++;  
    if (readcount == 1) wait(wrt);  
    signal(mutex);  
  
    // reading is performed  
  
    wait(mutex);  
    readcount--;  
    if (readcount == 0) signal(wrt);  
    signal(mutex);  
}
```





# The first readers-writers problem



# The second readers-writers problem

- The former solution is sub-optimal.
  - Because it is possible that a reader R1 might have the lock, a writer W be waiting for the lock, and then a reader R2 request access.
  - It would be foolish for R2 to jump in immediately, ahead of W; if that happened often enough, W would **starve**.
  - Instead, W should start as soon as possible.
- This is the motivation for the **second readers-writers problem**, in which the constraint is added that **no writer, once added to the queue, shall be kept waiting longer than absolutely necessary**.
- This is also called **writers-preference**.

# Readers-Writers Problem: Solution #2

- int readcount = 0,  
writecount = 0;
- semaphore x = 1, y = 1,  
wrt = 1, red = 1;
- The structure of a  
writer process


```
while (true) {  
    wait(y);  
    writecount++;  
    if (writecount == 1)  
        wait(red);  
    signal(y);  
    wait(wrt);  
    // writing is performed  
    signal(wrt);  
    wait(y);  
    writecount--;  
    if (writecount == 0)  
        signal(red);  
    signal(y);  
}
```

# Readers-Writers Problem: Solution #2



- The structure of a reader process

```
while (true) {  
    wait(red);  
    wait(x);  
    readcount++;  
    if (readcount == 1) wait(wrt);  
    signal (x);  
    signal(red);  
    // reading is performed  
    wait (x);  
    readcount--;  
    if (readcount == 0) signal(wrt);  
    signal(x);  
}
```




# The second readers-writers problem

```
wait (y);  
    writecount++;  
    if (writecount == 1) wait(red) ;  
signal (y);  
wait (wrt);  
.....
```

```
wait (red);  
.....
```



wait (red)



# The third readers-writers problem

- In fact, the solutions implied by both problem statements result in **starvation** — the first readers-writers problem may starve writers in the queue, and the second readers-writers problem may starve readers.
- Therefore, the **third readers-writers problem** is sometimes proposed, which adds the constraint that **no process/thread shall be allowed to starve**; that is, the operation of obtaining a lock on the shared data will always terminate in a bounded amount of time.
- Solutions to the third readers-writers problem will necessarily sometimes **require readers to wait even though the share is opened for reading**, and sometimes **require writers to wait longer than absolutely necessary**.

# Readers-Writers Problem: Solution #3




- `int readcount = 0;`
- semaphore `mutex = 1, wrt = 1, S=1;`
- The structure of a writer process

```
while (true) {  
    wait (S);  
    wait (wrt);  
    // writing is performed  
    signal (wrt);  
    signal (S);  
}
```

# Readers-Writers Problem: Solution #3



- The structure of a reader process



```
while (true) {  
    wait (S);  
    wait(mutex);  
    readcount++;  
    if (readcount == 1) wait(wrt);  
    signal(mutex);  
    signal (S);  
  
    // reading is performed  
  
    wait(mutex);  
    readcount--;  
    if (readcount == 0) signal(wrt);  
    signal(mutex);  
}
```


```
while (true) {  
    wait(mutex);  
    readcount++;  
    if (readcount == 1) wait(wrt);  
    signal(mutex);  
    // reading is performed  
  
    wait(mutex);  
    readcount--;  
    if (readcount == 0) signal(wrt);  
    signal(mutex);  
}
```

Solution#1

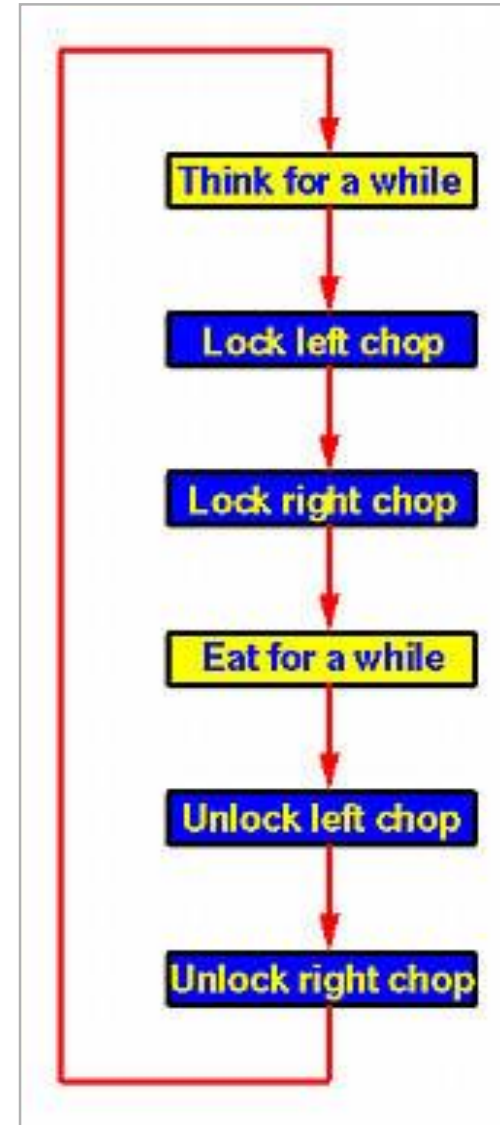




# Dining-Philosophers Problem

- Five philosophers, either **thinking** or **eating**
- To eat, two chopsticks are required
- Taking one chopstick at a time




# Dining-Philosophers Problem



# Dining-Philosophers Problem

- Shared data
  - Bowl of rice (data set)
  - Semaphore chopstick [5] initialized to 1



# Dining-Philosophers Problem

- The structure of Philosopher  $i$ :  

```
While (true) {  
    wait ( chopstick[i] );  
    wait ( chopstick[ (i + 1) % 5] );  
    // eat;  
    signal ( chopstick[i] );  
    signal ( chopstick[ (i + 1) % 5] );  
    // think;  
}
```
- Diagram illustrating the structure of Philosopher  $i$  and the corresponding chopstick operations:
- wait ( chopstick[i] );** → **get chopsticks** (left)
  - wait ( chopstick[ (i + 1) % 5] );** → **get chopsticks** (right)
  - signal ( chopstick[i] );** → **free chopsticks** (left)
  - signal ( chopstick[ (i + 1) % 5] );** → **free chopsticks** (right)

# Dining-Philosophers Problem

- Possible solutions to the deadlock problem
  - Allow at most **four** philosophers to be sitting simultaneously at the table.
  - Allow a philosopher to pick up her chopsticks only if **both chopsticks are available** (note that she must pick them up in a critical section).
  - Use an asymmetric solution; that is,
    - odd philosopher: **left first, and then right**
    - even philosopher: **right first, and then left**
- Besides deadlock, any satisfactory solution to the DPP must avoid the problem of **starvation**.

# Solution 2 for Dining-Philosophers Problem

- semaphore mutex=1;

```
void philosopher(int i) {  
    while(TRUE) {  
        think( );  
        P(mutex);  
        take_chopstick (i);  
        take_chopstick ((i + 1) % N);  
        eat( );  
        put_chopstick (i);  
        put_chopstick ((i + 1) % N);  
        V(mutex);  
    }  
}
```

# Solution 3 for Dining-Philosophers Problem

---

- S1 THINKING...
- S2 I am HUNGRY
- S3 If my left neighbor or my right neighbor is EATING then block myself; else goto S4
- S4 Pick up both chopsticks
- S5 EATING ...
- S6 Put down the chopsticks and wake up the left neighbor if he can EAT
- S7 Put down the chopsticks and wake up the right neighbor if he can EAT
- S8 Goto S1

# Solution 3 for Dining-Philosophers Problem

- Define the data structures:

```
#define N                5
#define LEFT             (i+N-1)%N
#define RIGHT            (i+1)%N
#define THINKING 0
#define HUNGRY          1
#define EATING           2
int  state[N];
semaphore  mutex;        // initial value 1
semaphore  s[N];         // initial value 0
```



# Solution 3 for Dining-Philosophers Problem

```
void philosopher(int i) // i: 0~N-1
{
    while(TRUE)
    {
        S1 ————— think( );

        S2-S4 ————— take_chopsticks(i);

        S5 ————— eat( );

        S6-S7 ————— put_chopsticks(i);
    }
}
```

# Solution 3 for Dining-Philosophers Problem



- // Pick up both chopsticks, or block

```
void take_chopsticks(int i) // i: 0~N-1
{
    P(mutex);
    state[i] = HUNGRY;
    test(i);
    V(mutex);
    P(s[i]);
}
```


```
void test (int i)
{
    if(state[i] == HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING )
    {
        state[i] = EATING;
        V(s[i]);
    }
}
```

# Solution 3 for Dining-Philosophers Problem

- // put down the two chopsticks and wake up the neighbors if necessary

```
void put_chopsticks(int i)
{
    P(mutex);
    state[i] = THINKING;
    test(LEFT);
    test(RIGHT);
    V(mutex);
}
```

# Sleeping Barber Problem



# Sleeping Barber Problem

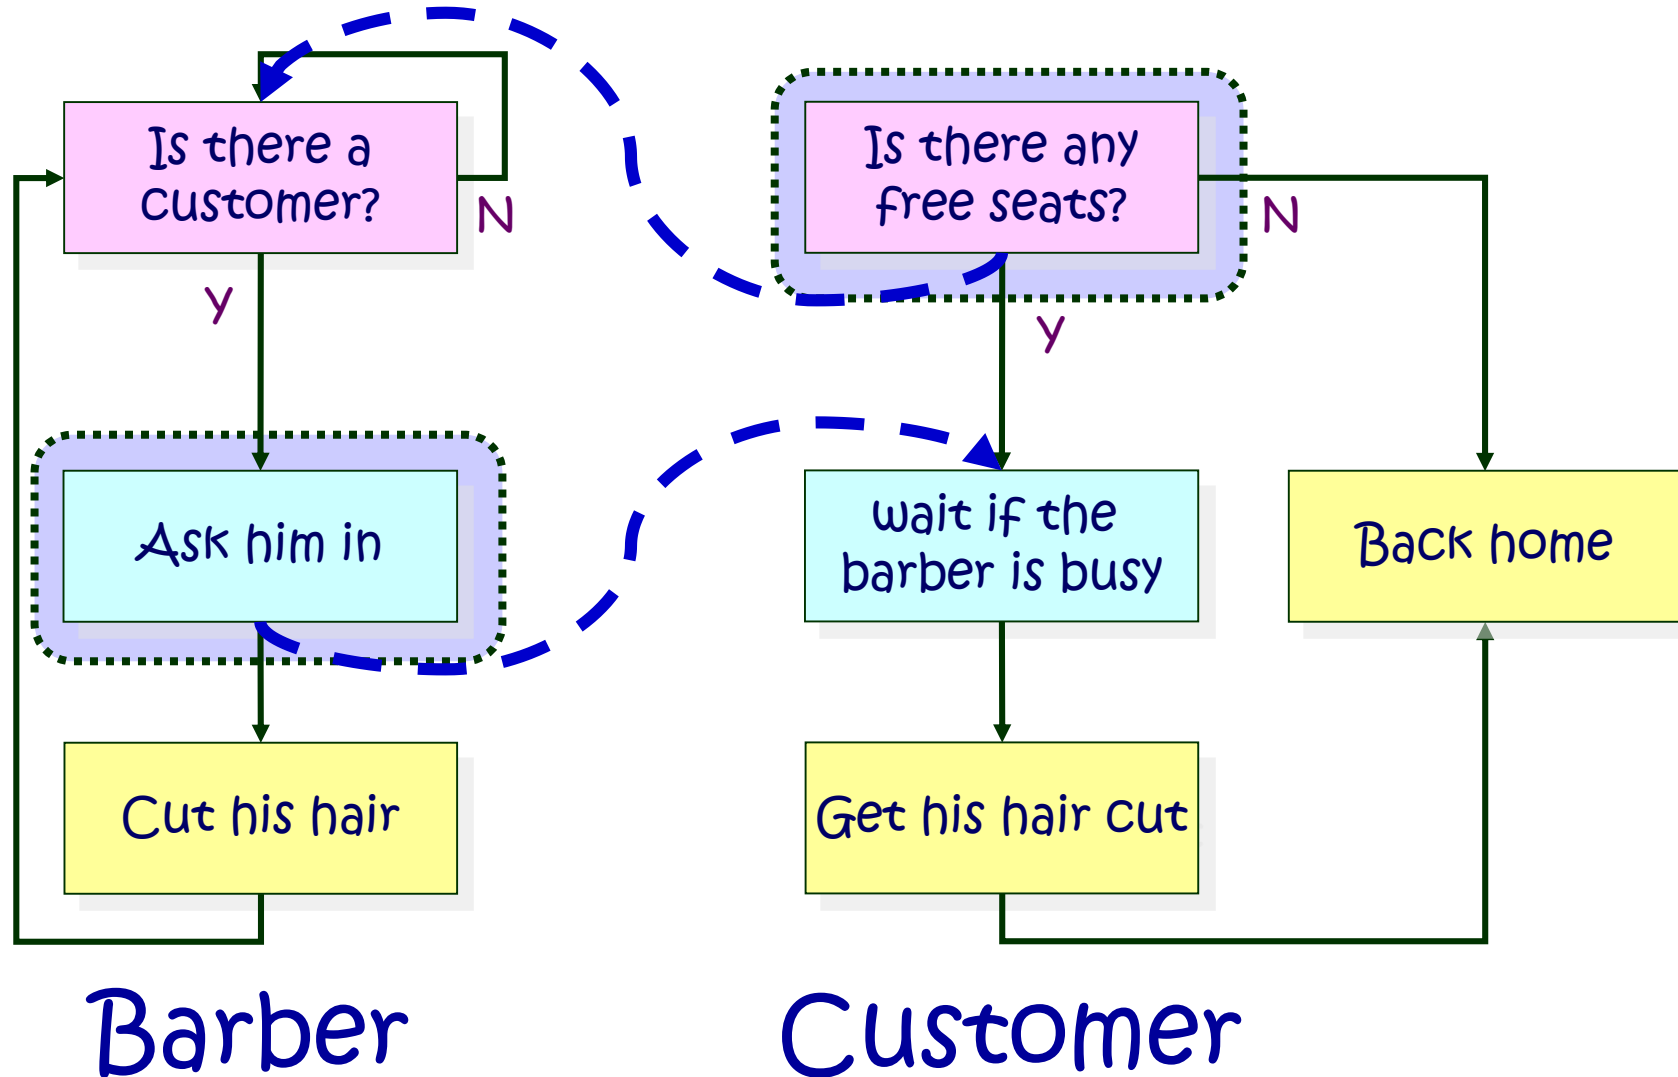
- The **sleeping barber problem** is a classic inter-process communication and synchronization problem between multiple operating system processes.
- There is a hypothetical barber shop with **one barber**. The barber has **one barber chair** and a waiting room with **a number of chairs** in it.



# Sleeping Barber Problem

- When the barber finishes cutting a customer's hair, he dismisses the customer and then goes to the waiting room to see if there are other customers waiting.
  - If yes, he brings one of them back to the chair and cuts his or her hair.
  - If no, he returns to his chair and sleeps in it.
- Each customer, when he arrives, looks to see what the barber is doing.
  - If the barber is sleeping, then he wakes him up and sits in the chair.
  - If the barber is cutting hair, then he goes to the waiting room.
    - If there is a free chair in the waiting room, he sits in it and waits his turn.
    - If there is no free chair, then the customer leaves.

# Sleeping Barber Problem



# Sleeping Barber Problem

- You need (as mentioned above):

Semaphore Customers = 0;    // #waiting customers

Semaphore Barber = 0;

Semaphore accessSeats (mutex) = 1;

int NumberOfFreeSeats = N    // total number of seats

Customers {  
    > 0, there are waiting customers  
    = 0, there are no waiting customers  
    = -1, the barber is sleeping

Barber {  
    = 1, the barber is not busy  
    = 0, the barber is busy cutting or sleeping  
    < 0, there are waiting customers



# The Barber Process

```
Void barber(void)
{
    while(true) {                // runs in an infinite loop
        P(Customers)             // tries to acquire a customer
                                // – if none is available he goes to sleep
        P(accessSeats)           // at this time he has been awakened
                                // - want to modify the number of available seats
        NumberOfFreeSeats++      // one chair gets free
        V(Barber)                 // the barber is ready to cut
        V(accessSeats)            // we don't need the lock on the chairs anymore
        cut_hair();               //here the barber is cutting hair
    }
}
```

# The Customer Process

```
void customer(void)
{
    while(true) {                // runs in an infinite loop
        P(accessSeats)           // tries to get access to the chairs
        if ( NumberOfFreeSeats > 0 ) { // if there are any free seats
            NumberOfFreeSeats--    // sitting down on a chair
            V(Customers)           // notify the barber, who's waiting that there is a customer
            V(accessSeats)         // don't need to lock the chairs anymore
            P(Barber)              // now it's this customer's turn, but wait if the barber is busy
            //here the customer is having his hair cut
        } else {                 // there are no free seat, tough luck
            V(accessSeats)         // but don't forget to release the lock on the seats
            //customer leaves without a haircut
        }
    }
}
```